

A BUILD MEMOIR · VERIFY EVERYTHING

The Honest Machine

*How one person and an AI built a
complete telecom suite — from catalog to
cash to campaigns to copilots — and
proved every piece of it*

Built with **Claude Fable 5** · genalpha-bss · thirty-three components, six channels, and a discipline

PREFACE

What this book is

This is the story of building a telecom Business Support System — the software an operator sells, bills, and serves customers through — from an empty repository to thirty-three composable components, six customer surfaces, and fifty-eight end-to-end test suites that all pass in a real browser. It was built by one person working with an AI pair, over a compressed and relentless arc, held together by a single discipline: *nothing is real until it is proven, end to end, against the real thing.*

It is a *memoir*, not a manual. There is a reference book to be written about how the pieces fit; this is not it. This is about the decisions, the dead ends, the two-in-the-morning debugging, and the method — and about what it means that a system this size was built the way it was.

Three convictions run through every chapter. The first is **vendor-neutrality**: any identity provider, any database, any message broker, any AI model, any payment processor, any number-porting authority. Nothing operator-specific is hardcoded, ever. The second is **verification**: official conformance kits, real-database migration tests, isolation proofs, and every feature driven through a real browser. The third is **honesty about what's a stand-in**: a thin network layer, a mock payment processor, a shaped-but-unconnected clearing house — each named plainly, so a demo stays credible instead of turning naive.

And beneath all three, the thread that makes this book unusual: it was built with an AI as the executing partner — Anthropic's **Claude Fable 5**, the most capable generally-available model of its moment. Not autocomplete. A collaborator that wrote the services, argued the architecture, made real mistakes, and got caught by the tests. The book is, quietly, a field report on what building at this scale with a model this capable actually feels like.

On the jargon. This is a book about a TM Forum ODA BSS, so it can't avoid the vocabulary. The rule is that it never assumes you already know a term: each chapter explains the one or two standards it needs as the story reaches them, and the appendix collects the full "what every TMF component means" reference. You should be able to read this having never heard of TM Forum and come out fluent.

Every claim in these pages has a commit, a test, or a screenshot behind it. That is not a stylistic choice. It is the whole argument.

ACT ONE

The First Cut

*Why build a Business Support System at all — and
the discipline that would define everything that
came after.*

The lock-in that started it

"The most dangerous phrase in the language is: we've always done it this way."

— Grace Hopper

Every telecom operator on earth runs software to answer three questions: what can a customer buy, what do they owe, and how do we serve them when something breaks. That software is called the BSS — the Business Support System — and it is, almost universally, the place operators are most trapped. The catalog, the ordering, the billing, the care: bought as one enormous suite from one enormous vendor, welded together over a decade, impossible to leave. The industry has a word for the escape hatch it wishes it had — *composable* — and a standards body, TM Forum, that publishes the blueprint for it. What almost nobody has is the thing built end to end, out in the open, where you can watch it work.

So that was the bet. Build the composable BSS the industry keeps promising, as a real product, and make it *vendor-neutral* from the first line of code. The operator I know best runs one particular identity stack; the product must never depend on it. It must run against any OpenID Connect provider, any PostgreSQL, any Kafka-protocol broker. Neutrality wasn't a feature to add later. It was the shape of the thing.

The first five components were the spine of any BSS. A **product catalog** (the shelf: what exists to sell, at what price, in what bundles). **Product ordering** (the checkout and its aftermath). **Product inventory** (what each customer actually has). A **party** component (the people and companies you do business with, kept deliberately separate from their login). And a **gateway** — one front door, so the outside world touches exactly one thing. Each spoke a published TM Forum Open API, which meant that from day one the system talked in a language other vendors' software already understood.

Constraints chosen early are cheap on day one and priceless on day one hundred.

That is the whole thesis of the first act, and it would be tested again and again. Every time the temptation appeared to hardcode something — the IdP, the number format, the country — the neutrality constraint said no, build a seam. Seams cost a little now. But a system made of seams is a system you can take apart, swap pieces of, and sell to an operator who already owns half of it. A monolith is a system you can only sell whole, to someone who owns none of it. There are not many of those customers left.

The lesson. The expensive decisions in software are the early, invisible ones. Choosing "any IdP, any database, any broker" on the first day looks like over-engineering. It is the opposite: it is the only cheap moment to make that choice.

Proof, or it didn't happen

"Simplicity is prerequisite for reliability."

— Edsger W. Dijkstra

The second decision mattered more than the first, though it looked like a chore. How do you *know* a thing works? The easy answer — write a unit test, watch it go green — is a comfortable lie. A unit test proves that the code does what the code's author thought it should. It says nothing about whether the author understood the problem, whether the pieces fit, or whether a real customer clicking a real button gets a real result.

So the answer became doctrine, and the doctrine was severe. First: the original components would pass TM Forum's official **Conformance Test Kits** — an independent, published suite that proves a component really implements a standard rather than merely resembling it. Not "we think it's TMF620." An outside authority agrees it's TMF620. Five kits at first; the doctrine would compound until, near the end of this story, *eleven* kits agreed — and the two that deliberately didn't became one of the book's most useful lessons. Second: every feature, without exception, would be driven through a *real browser* — register, configure, add to cart, pay, watch the order complete, watch the notification arrive. If a human couldn't do it in a browser, it wasn't done.

There was a subtler third rule, learned the hard way. The tests ran against an in-memory database that imitated PostgreSQL. Close, but not identical — and "close but not identical" is precisely where the migrations break. So the migrations grew their own test that spun up a *real* PostgreSQL and ran every schema change against it. The in-memory database is for speed. The real one is for truth.

Verification isn't overhead. It is the thing that lets you move fast and still trust the result.

This is the discipline that would, later, make it rational to build alongside an AI at a speed no solo human could sustain. When the model wrote a service, the tests didn't care that it was confident. They cared whether an order actually placed, whether a row was actually isolated, whether a real browser actually saw the result. The first "ALL CHECKS PASSED" scroll past the terminal was not a milestone. It was the establishment of a court of law that would preside over everything that followed.

The lesson. Choose your arbiter before you need it. A test that drives the real system through the real interface is a fact. Everything else — the author's confidence, the model's fluency, the green unit test — is an opinion.

The pair

"Programs must be written for people to read, and only incidentally for machines to execute."

— Harold Abelson & Gerald Jay Sussman

The collaborator has a name: **Claude Fable 5**. It wrote most of the lines in this system. Describing the working method is the honest heart of the book, because the method is the thing that's actually new. A human held intent and judgment — what to build, whether it was enough, which architectural fork to take. Fable 5 held execution and breadth — the services, the boilerplate, the cross-component wiring, and the ability to keep an entire thirty-component system in view at once, which no human can.

What it felt like, most of the time, was uncanny. You describe a component — "a promotion service, anonymous code validation, redemption tied to an order, a discount line on the bill" — and a few minutes later it exists, wired into the gateway, event-publishing, tenant-isolated, with a test that drives it through the browser. The model doesn't just write the service; it remembers the seventeen other services it has to speak to, and the tenant rules, and the event envelope, and it wires them without being told.

But a memoir that only showed the magic would be the same lie as a passing unit test. Fable 5 was wrong, sometimes, and wrong in the specific way capable models are wrong: fluently. It misdiagnosed a gnarly network bug twice before the evidence forced the truth. It once added a configuration value to the wrong service because a text-replacement matched the first occurrence instead of the intended one — a mistake no human would make and no human would catch by reading, but the test caught it in seconds, because the number never activated on the ported line.

The verification discipline isn't a brake on a capable model. It is the thing that lets you use the capability at full speed without fooling yourself.

That is the shape of the partnership, and it will recur in every chapter that follows: the model proposes at a breadth and speed that is genuinely startling; the human decides and, crucially, the *tests decide too* — not the model's confidence, and not the human's. Every commit in the repository is co-authored by Claude Fable 5. The credit is on the record, not just in the prose.

The lesson. AI-assisted building is only as trustworthy as the verification around it. Give the best model the whole system to hold, and a discipline that catches it when it's wrong, and the pair can build what used to take a team — provided nobody ever mistakes fluency for correctness.

ACT TWO

Scale & the Multitenant Reckoning

*Five components became thirty. One deployment
learned to hold two operators who must never see
each other. And the whole invisible machine learned
to show its face.*


One by one

"A distributed system is one in which the failure of a computer you didn't know existed can render your own computer unusable."

— Leslie Lamport

After the first five, the components came like tide. Stock, so the shop could tell the truth about the last handset. Payment, so money could move. Billing, so a customer could be charged for what they used. Qualification, so an address could be asked whether fibre even reached it. Appointments, tickets, interactions, communications, carts, usage, agreements, promotions, roles, addresses, recommendations, the card vault, documents. Each one a small, self-contained business, each one speaking a published standard, each one wired into the same nervous system.

That nervous system was the thing that made thirty components a system instead of thirty programs. The rule was simple and never broken: a component never reaches into another's database, and never calls another and waits. It writes what happened to its own tables, in the same transaction as its own work, and a background process ships that record onto a shared log. Other components listen. When an order completes, the ordering component doesn't *tell* billing, inventory, and communications to act. It simply records, in its own words, that an order completed — and the others, listening, do their part.



The transactional outbox: write-what-happened, ship-it-to-the-log, let listeners react. Loose coupling as a physical property of the code, not a diagram.

Why go to this much trouble? Because the alternative — component A calls component B and waits — is the thing Lamport warns about. Chain enough synchronous calls together and a hiccup in a component you'd forgotten existed freezes a checkout. Record-and-react has the opposite failure mode: if billing is briefly down, the events wait patiently on the log, and billing catches up when it wakes. The customer's order completed the moment the order completed. Everything else is downstream, and downstream can be a little late without anyone noticing.

Thirty programs become a system when they stop calling each other and start telling the truth about what just happened.

The lesson. The coupling in a system is not in its diagram; it's in its call graph. Components that record facts and react to facts are loosely coupled in the only way that matters — the way that survives one of them falling over at 3 a.m.

Two operators, one deployment


Data isolation is not a feature you add. It is a property you can lose — silently, on any query, forever.

— from the build log

Here is the reckoning the second act is named for. A product you can sell to many operators is a product that runs many operators *at once*, on one deployment, sharing the same tables — and it must be impossible, not merely unlikely, for one operator to see another's customers. Get this wrong and there is no apology that suffices. This is the chapter where the whole architecture had to earn its trust.

The first question was: what is a tenant? The tempting answer is "a row in a tenants table," but that just moves the problem. The answer chosen was harder and truer: a tenant is a *verified identity provider*. If you can prove you were authenticated by the issuer that operator A registered, you are operator A's user — no separate account system, no shadow directory, no operator-specific code. Bring your own OpenID Connect issuer and you bring your own tenancy. That's neutrality and multitenancy solved by the same stroke.

But identity at the door is not isolation at the database. The load-bearing wall is *Postgres row-level security*. Every tenant-owned table carries a `tenant_id`, and the database itself — not the application, which can have bugs — refuses to return a row whose tenant doesn't match the one stamped on the current connection. The application runs as a restricted role that *cannot* see across tenants even if a query forgets its `WHERE` clause. There is a narrow, audited escape hatch for genuine system work, and migrations run as the owner. Everything else lives inside the wall.



Three layers, and only the last one is trusted. Identity resolves the tenant; the connection carries it; the database enforces it. A forgotten WHERE clause is caught by the wall, not by hope.

And then — because this is a book about proof — the isolation got its own test. Not a comment swearing it was safe. A test that logs in as operator A, creates a customer, logs in as operator B, and demands that B cannot see it — against a real PostgreSQL with the real policies. The gateway got its own trick too: an anonymous visitor typing a shop's hostname is routed to the right tenant by the hostname alone, so a customer who hasn't logged in yet still lands in the right operator's world.

A promise of isolation is worth exactly the test that tries to break it and fails.

The lesson. Put the guarantee at the lowest layer that can enforce it. Application code that "always filters by tenant" is one forgotten clause from a breach. A database that physically cannot return the wrong row is a different kind of promise — the kind you can sell.

War stories of scale

"Plan to throw one away; you will, anyhow."

— Frederick P. Brooks Jr.

Thirty services do not fit on a laptop by wishing. This is the chapter of the unglamorous fights — the ones that never appear in an architecture diagram and consume whole evenings. The container base image that had no build for the laptop's chip, so every image had to be pinned to one that did. The memory: thirty Java processes, each wanting a comfortable heap, summing to more RAM than the machine had, until every service was told firmly how little it was allowed. The builds parallelised until the fan roared, then throttled to four at a time so the machine stayed usable.

The test harness had its own private hell. It spun up real databases in throwaway containers — the right instinct, the one that catches migration bugs — but it talked to the container engine over a socket that, on this particular setup, lived somewhere non-standard, and it guessed an API version the engine didn't speak. Two environment variables, discovered over an hour that felt like three, and it worked. The kind of fix that is one line in the end and a small saga in the getting-there.

None of this is architecture. All of it is the tax you pay to run the architecture, and a memoir that skipped it would be selling the same fantasy as a demo that only ever runs on the presenter's machine. The reason to write these down is not nostalgia. It's that the next person — or the next model — hits the same walls, and a paragraph here saves them the evening.

The distance between "it works" and "it works on a machine that isn't yours" is where most software quietly dies.

The lesson. The environment is part of the system. A build that only succeeds on one laptop, with one chip, one socket path, one memory profile, is not done — it's performing. Write the walls down; they are load-bearing knowledge.


Making it visible

"Perfection is achieved, not when there is nothing more to add, but when there is nothing left to take away."

— Antoine de Saint-Exupéry


A customer never sees a component. They see a shop. For twenty-odd components, the system had been real but faceless — provably correct, and invisible. This is the chapter where it grew four faces: a **storefront** where a customer browses and buys; a **CSR console** where an agent serves them; an **admin console** — the back office — where an operator runs their catalog, campaigns, and staff; and a **mobile app** in the customer's pocket.

The subtle demand was that the faces be the operator's, not mine. A tenant is a brand, and a brand has a colour. So the channels learned to wear the tenant: the same storefront, rendered in one operator's identity and then another's, from nothing but their configured brand and logo. The two shops below are the *same code*, the same components, the same database — wearing two different operators.



One storefront, two operators. Same components, same deployment, different brand — white-labeling as a runtime property of the tenant, not a fork of the code.

The mobile app closed the loop: the same BSS, the same tenant rules, in the customer's hand.



The fourth channel — the account in your pocket, served by the same thirty components.

Making the system visible changed what it *was*. A correct-but-invisible platform is a thing you have to be an engineer to believe in. A shop you can click through, in an operator's own colours, is a thing you can show a buyer. The act closes with the machine no longer merely proven, but legible.

The lesson. The channel is where correctness meets belief. All the conformance in the world persuades no one who can't see it; a shop that wears the customer's brand and just works is the argument the tests were always making, finally visible.

ACT THREE

The Intelligence Turn

*Adding AI to a BSS without surrendering to a single
vendor, without leaking a customer's secrets, and
without pretending a rule is a brain.*

AI on our terms

"All models are wrong, but some are useful."

— George E. P. Box

The moment you add AI to a product, a gravity well opens: one vendor's model, one vendor's API, one vendor's prices, one vendor's outage becoming yours. For a platform whose entire conviction is neutrality, capitulating there would have been a betrayal of the first chapter. So the AI arrived the same way everything else had — behind a seam. An *intelligence* component with one job: to be the only place in the system that talks to a language model, and to talk to *any* of them.

A stub model for tests that need no cloud. An OpenAI-compatible adapter for the dozens of providers that speak that dialect. An Anthropic adapter for Claude. Each selected by a flag, per tenant, so operator A can run one model and operator B another, on the same deployment. The model is a setting, not an architecture.

Two rules were welded on before the first feature shipped, because they are the two ways AI quietly poisons a BSS. The first: **always-on redaction**. Before any customer text reaches any model, emails and phone numbers are stripped to [email] and [phone]. Not optionally. Not configurably. Always. The second: **an audit ledger**. Every prompt and every completion is written to a per-tenant, isolated table, so an operator can answer the question every regulator will eventually ask — *what did the AI see, and what did it say?*

A model you can swap is a model that can't own you. A prompt you can audit is a model you can defend.

This chapter is where the first act's conviction met the decade's most seductive lock-in and did not blink. Everything intelligent that follows — the copy assistant, the copilot, the churn engine — lives behind this seam, redacted and logged, portable across models by a single flag.

The lesson. Adopt the capability, refuse the capture. An AI seam with redaction and audit welded in is how you get the upside of a language model without betting your product, your customers' privacy, or your compliance story on one vendor's roadmap.

Where AI belongs

The right question is never "can AI do this?" It is "should a human be reading a draft, or reading a decision?"

— from the build log

A seam is a capability, not a use. The harder question was where in a BSS an AI actually earns its place — and the answer that held up was consistent: *AI drafts; humans decide*. It writes the first version; a person ships it. Every feature was chosen to sit on the draft side of that line.

The first was a campaign copy assistant. A marketer building a customer journey types a rough intent and gets back a subject line and a body, on-brand, with the promo code guaranteed to appear — but never inventing the discount amount, a bug caught early when a model cheerfully promised "20% off" for a code worth ten. The marketer edits and sends. The AI saved the blank page, not the judgment.

GenAlpha

BACK OFFICE

demo

Sign out

Product Offerings

Product Specifications

Product Offering Prices

Product Stock

Customer Bills

Serviceable Areas

Appointments

Campaigns

AI Audit

Campaigns

36 total

NEW

Name *

Summer win-back

Trigger *

Churn risk detected (AI scorer)

State filter

e.g. completed (optional)

Promo code to include

—

Message subject *

Message — {code} inserts the promo code *

AI assist

warm win-back for customers whose plan is ending, offer the code

Draft

Create

NAME	STATUS	TRIGGEREVENTTYPE	PROMOTIONCODE	REACHED	
Smoke welcome	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783787634183	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783787720216	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783789023110	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume

The admin console's campaign builder. The AI drafts the subject and body; the marketer owns what actually goes out.

The second went to the support desk: a CSR copilot. When an agent opens a customer, the copilot reads the 360 — the products, the tickets, the timeline — and offers a two-line summary and a suggested reply. Crucially, it holds no credentials of its own; it sees exactly what the agent is allowed to see, through the same tenant wall, and not one row more. The agent reads a draft, not a verdict.

GenAlpha

CSR CONSOLE

GENALPHA-RETAIL

Customers

Tickets

Stock

AA Anna Agent

Sign out

GG

Gina Guest

013f3d14-b1ba-43a7-b978-2e0071df3f97 · Gästgatan 7, 22233 Göteborg

The provided customer data is summarized deterministically (stub provider — configure a real model for genuine insight). Nothing here left the machine.

Review the open items in the 360 with the customer.

Check whether the current plan still fits their usage.

Drafted by stub (deterministic-template) — verify before acting.

Orders

GenAlpha One Home & Mobile

Jul 11, 12:34 AM

acknowledged

Complete

Cancel

Services

Nothing provisioned.

Usage this month

No usage recorded.

Agreements

No agreements.

Bills

No bills.

Appointments

Installation: GenAlpha One Home & Mobile

Jul 14, 01:00 PM

cancelled

Active cart

The customer's cart, live — assisted checkout starts here.

Promotions & payment

No promotions or saved cards.

Suggest next

GenAlpha One Home & Mobile

#1

GenAlpha TV Max

#2

Samsung Galaxy S26

#3

Tickets

No tickets.

Raise a ticket for this customer...

Raise ticket

Interactions

No interactions logged.

The CSR copilot: a customer summary and a suggested reply, drawn only from what this agent may already see. It drafts; the agent decides.

The discipline is quiet but total. Nowhere does the AI place an order, issue a refund, or close a ticket on its own. It compresses, it suggests, it drafts. The moment of consequence stays human — which is exactly why an operator can turn it on without fear.

The lesson. The safe and valuable place for AI in a system of record is the draft, not the decision. Put it where a human reads its output before anything irreversible happens, and you get the speed without the liability.

The engine that learns

A rule you can read is honest about being a guess. A model that learns is honest about being wrong until it isn't.

— from the build log

Churn — a customer about to leave — is the prediction every operator wants and few trust, because most churn scores are black boxes that can't say *why*. So the churn engine was built to be honest twice over.

It began as rules a human can read and argue with: a commitment ending within the month, an allowance running consistently against its ceiling, a cluster of support tickets, a service degraded during an outage. Four plain signals, each defensible, each explainable to the customer it flags.

But rules are a floor, not a ceiling, and the operator's real ask was pointed: *"I want it to learn once it's live, and become production-quality in a few months — or be trained on old data."* So underneath the rules sits a real, if modest, learner. Every day it snapshots each customer's features; when customers actually leave — or actually port their number out, the strongest signal of all — it labels the outcome. From those snapshots and outcomes it fits a per-tenant model that sharpens over time. Cold-start on rules; warm up on reality; import history if you have it.

Start with a rule you can defend. Earn the right to a model you can trust.

And because it lives behind the same event backbone as everything else, a churn signal isn't a number on a dashboard nobody reads. When the engine flags a customer, it emits an event — and the campaign component, listening, can reach out before the customer reaches for the exit. The prediction becomes an action without a human relaying it. That closed loop — detect, decide, act — is the whole point, and it is built from the same record-and-react nervous system as an order completing.

The lesson. Don't ship a black box into a domain that demands explanations. Ship the readable rules first, let them earn trust, and grow the learned model underneath where the stakes reward it — with the outcomes reality hands you as the labels.

ACT FOUR

Autonomy Accelerated

*A stadium, a 5G slice, edge GPUs sold by the token,
and a network that heals itself while a B2B deal
closes itself. The Catalyst, rebuilt on our own BSS.*

The Catalyst

"The best way to predict the future is to invent it."

— Alan Kay

There was a proof-of-concept I'd seen at an industry event and helped write up: selling the network edge as an AI platform. A football stadium wants a private 5G slice with ultra-low latency and on-site AI capacity for real-time video. A business salesperson uses an AI assistant to shape the lead; the operator's systems autonomously judge whether the network can even do it, propose an edge-AI upsell, price it, provision it — and then, when a fibre is cut mid-event, heal around the damage and keep the promise. It was a vision demo, stitched from many vendors' parts.

The dare I set myself was to rebuild it — not as a slideshow, but running, on the very BSS this book has been building. Every glossy moment of that vision would have to become a real request against a real component that really did the thing. If the BSS was what I claimed, it should be able to *host* the future, not just narrate it.

A vision demo shows you the destination. The dare was to lay the track.

This act is that rebuild, moment by moment: intent, feasibility, the upsell, token pricing, self-healing, and — the strangest and most modern part — one company's AI agent negotiating with another's, machine to machine. Each one grounded in a component you've already met, extended just far enough to carry the weight.

The lesson. The test of a platform is whether someone else's flagship demo can be rebuilt on top of it as working software. A BSS that can host the industry's most ambitious vision, honestly, is a BSS that was built right.


Intent, feasibility, and the honest "no"

Tell the network what you want, not how to build it — and let it tell you the truth about physics.

— from the build log

The old way to sell a network service is a form with forty fields. The new way, the one the standards call *intent*, is a sentence: "a sub-20-millisecond slice at the stadium, with AI inference capacity." The customer states the *what*. The system works out the *how* — or works out that there is no how.

That last part is where most autonomy demos cheat, and where this one refused to. The orchestrator runs a feasibility check against a hard, physical rule: a sub-20ms guarantee requires an edge-GPU pool actually covering the place in question. If one exists, the intent is feasible and the system proposes the slice. If none does, the answer is **INFEASIBLE** — and the message says, in as many words, that this is *physics, not policy*. It will not promise latency that light itself can't deliver.



Autonomous feasibility with a spine. When the physics allows it, the system proposes — and upsells edge-AI. When it doesn't, it says so plainly, and doesn't sell a lie.

When it is feasible, the same autonomy that could have said no instead says "yes, and": a stadium buying low latency for video is a stadium that wants AI inference at the edge, so the system proposes it unprompted. The upsell isn't a hardcoded banner; it's a consequence the orchestrator reasons its way to. Autonomy that can refuse is autonomy you can believe when it recommends.

The lesson. An autonomous system that can only say yes is a liar with a schedule. The credibility of every "yes, and here's the upsell" rests entirely on its willingness to say "no, and here's why" when the physics demands it.

Priced by the token

The unit of billing tells you what business you're really in.

— from the build log

Here the BSS earned its keep in a way no slideshow can. An edge-AI service isn't sold by the month; it's sold by the inference — by the *token*, the same unit the AI industry itself runs on. So the catalog gained an Edge AI Inferencing offering with real token metering: fifty million inferences included, then a per-million overage price, defined the standard way and rated the standard way.

This is where the earlier, unglamorous components suddenly paid off. Usage metering — built chapters ago for gigabytes and minutes — didn't care that the unit was now AI tokens. Billing didn't care. The commitment, the overage, the line on the invoice: all of it already existed, because it had been built as a general capability rather than a special case. Selling AI by the token required almost no new billing code, because the billing had been built right the first time.

The best feature is the one you already built for another reason and discover you can charge for.

And the commercial motion around it was the intent loop made concrete: the intent produced a priced **quote** — born from the conversation, itemising the slice and the token bundle, wrapped in an AI-written narrative explaining what the customer was buying and why. Accept the quote and it became an order. Lead to intent to quote to order, most of it autonomous, all of it against real components.

The lesson. Build capabilities, not features. A usage-and-billing engine built generically will one day meter a unit you hadn't imagined — AI tokens — and charge for it without a rewrite. That optionality is the compounding return on doing it properly.

The network that heals itself

Resilience is not the absence of failure. It is the presence of a plan that runs without you.

— from the build log

The dramatic beat of the original vision was a fibre cut, mid-event, and a network that healed around it before anyone in the stadium noticed. It's the moment that separates a brochure from a system, because healing is easy to *say* and unforgiving to *do*. So it got built as an actual loop, wired to the assurance component that had been watching the network all along.

A critical alarm arrives on the fibre route. Assurance, which turns alarms into tracked service problems, raises one. The self-heal service sees it and acts: it re-homes the affected service off the cut fibre and onto an edge GPU site that can still carry it, records the migration, opens a ticket so a human has an audit trail, and — once the service is stable on its new path — resolves the problem automatically. The SLA holds. The stadium keeps its slice. And every step left a record, because a self-healing system that heals invisibly is indistinguishable from one that's lying.



Every step recorded. A heal you can't audit is a claim, not a system.

The self-heal loop: alarm to problem to re-home to restored SLA, autonomous end to end, with a ticket left behind so a human can trace exactly what the machine did.

Autonomy you can trust is autonomy that shows its work.

The lesson. A self-healing system must be an *auditable* self-healing system. The re-home is the clever part; the ticket it opens on the way is the part that makes an operator willing to let it run unattended.

Agents talking to agents

The next integration boundary isn't an API. It's a conversation between two machines that each brought a model.

— from the build log

The strangest and most forward-looking beat came last. In the original vision, the buyer's side and the seller's side each had their own AI, and the deal moved between them machine-to-machine. That sounds like science fiction until you build it, and then it's just a protocol.

So the BSS grew a small server speaking the emerging standard for exactly this — a way for an *external* AI agent to drive the lead-to-order motion through a handful of well-defined tools: draft an intent, submit it, create a quote, accept a quote, list quotes. A salesperson's Claude, sitting entirely outside the BSS, could now shape a stadium deal by calling these tools — and on the other side, the operator's autonomous orchestrator answered with feasibility, an upsell, and a price. Two AIs, two companies, one deal, no human typing into a form.

What made this more than a party trick was that the tools weren't a special AI-only backdoor. They were thin wrappers over the exact same components a human channel used — the same intent, the same quote, the same order, the same tenant wall. The agent got no special powers and no special access. It simply drove the front door with its hands instead of a mouse.

When the integration layer becomes a conversation, the thing you must get right isn't the AI. It's that the AI touches nothing a human couldn't.

The lesson. Machine-to-machine autonomy is coming through the tool interface, and the discipline that makes it safe is the oldest one in this book: the agent acts through the same guarded front door as everyone else, with the same tenancy, the same limits, the same audit. No backdoors for robots.

ACT FIVE

Making It Real, Making It Legible

*Hardening the money and the identity, keeping and
leaving a phone number, running the whole thing
on a real cluster, and finally making the invisible
machine tell its own story.*

The money and the identity

Two things a telecom must never get wrong: charging a customer twice, and believing someone is who they aren't.

— from the build log

"Let's harden the BSS first" was the instruction, and it was the right one. A demo can be forgiven a hand-wave at payment. A product cannot. So the payment path was rebuilt to the standard a real processor demands. **Idempotency** first: every payment carries a correlator, so a customer who double-clicks, or a network that retries, gets the *same* payment — never a second charge. Then the full life of money: authorize a hold, capture it, void or refund it, each a real movement recorded with the processor's own reference.

And because real payments sometimes demand more — a bank insisting on a second factor — the path learned to pause. When the processor requires strong customer authentication, the API doesn't fail; it returns a "we need more" with the URL to complete it, and resumes when the customer has. The mock processor even carries test cards that decline or demand that step-up, so the hard paths get exercised on every run, not just the happy one. A real Stripe adapter sits behind the same seam, one flag away — because of course it does.

Identity got the same seriousness. Some purchases — a high-value plan, a regulated product — should require a *verified* identity, the kind a national eID like BankID provides. So the catalog can mark an offering as requiring verified identity, and ordering enforces it: no verified identity claim in your token, no order — a clear refusal with a header telling the channel to send you to step up. The verification itself is the identity provider's job, which is exactly where a neutral platform wants it.

Hardening is where a demo becomes a product: the same flow, now unable to charge you twice or take your word for who you are.

The lesson. The unglamorous guarantees — idempotent payments, real authorize/capture/refund, verified identity where the stakes demand it — are the line between something that demos and something you can put real customers and real money through. Harden before you scale.


Keep your number, or leave with it

The truest test of an open platform is whether a customer can walk away — and take their number with them.

— from the build log

Number portability is the quiet backbone of telecom competition. A customer will only switch operators if they can keep their number, and regulators mandate it precisely because it keeps operators honest. Building it meant confronting a truth the first chapter predicted: portability is *country-specific*. In Norway, a national clearing house called NRDB, under the regulator Nkom, brokers every port. Elsewhere it's a different authority with different rules, different formats, different cutover windows.

So porting got the same treatment as everything else: a seam. A porting gateway with a country-aware implementation — NRDB for Norway, shaped exactly as the real thing but not wired to it, and honestly labelled so — and per-country rules for number format, cutover timing, and regulator. "Vendor-neutral" had quietly become "regulator-neutral." The same code ports a number in Norway, Sweden, Britain, or the United States by consulting the rules for the country, not by branching on a hardcoded assumption.



Both directions of the door. Port-in draws the customer's real number instead of a pool number and holds the line until cutover. Port-out ceases the service, releases the number — and feeds the churn engine the truest label it will ever get.

Keeping your number when you arrive meant the orchestrator had to wait: instead of grabbing a fresh number from the pool, it holds for the port-in's cutover and then carries the customer's own number. Leaving with it meant building something the BSS didn't yet have — a service *cease* flow — because a port-out is also a termination. And a termination, the churn engine noticed, is the single most honest churn label in existence: not a prediction of leaving, but the fact of it, fed straight back into the model that tries to see the next one coming.

The lesson. The features that let a customer leave are the features that prove a platform is open. Build the exit as carefully as the entrance — country-aware, honest about stand-ins — and the same event that loses a customer teaches the system to keep the next one.

Running it for real

"In theory there is no difference between theory and practice. In practice there is."

— attributed to Benjamin Brewster

Everything so far had run on a laptop, in a chain of containers. That proves the software works; it doesn't prove the software *deploys*. Between the two lies a canyon full of the things that never show up until you cross it — resource requests that were fine on a laptop and starve a scheduler, config that was an environment variable and now must be a mounted secret, a component that came up instantly and now waits on ten others.

So the whole system was packaged the way it would actually ship — a Helm chart for Kubernetes, with infrastructure defined as code for real clouds — and then *soaked*: brought up on a real cluster and left to run, watched, prodded, and made to prove it could stand on its own for more than the length of a demo. A soak isn't a test that passes or fails in a second. It's the long, boring, essential question — does it stay up? — asked of the parts that only misbehave over time.

"It runs on my machine" and "it runs" are separated by a canyon, and the canyon is where products are made.

The soak had its own runbook and its own honesty: which components are core and which can be scaled back to fit a modest cluster, what to watch, what "healthy" means over hours rather than seconds. It is the least cinematic chapter in the book and one of the most important, because it is the difference between something that impressed a room and something an operator could actually run.

The lesson. Deployment is a feature, and soak is its test. A system that has never run unattended on a real cluster for hours has not been proven to run at all — it has been proven to *demo*, which is a different and smaller claim.

The invisible, made visible


If a business process runs and no one can watch it, you have built a machine of faith.

— from the build log

All those events, flowing between all those components, were the true life of the system — and completely invisible. You had to be an engineer reading logs to believe any of it was happening. The last component set out to fix that: **Live Flow**, a face for the nervous system itself.



The first attempt was an honest engineer's instinct — a graph of components with lines lighting up as messages passed between them. It was correct and it was useless to anyone who wasn't me. The feedback was exact and it reframed everything: *show it like a business-process engine shows a process. Like a layman can understand — a customer ordering, the order component activating, an order-created event, captured by something.* Not a wiring diagram. A story.

So Live Flow was rebuilt as a set of business processes told in plain English — Order to Activation, Lead to Assurance, Churn to Retention, Port-in to Keep-Your-Number, Port-out to Leave — each a row of labelled stages that *light up in order* as the real events arrive, teal for a normal step, violet when AI touched it, rose when something went wrong, fading as they cool.



Live Flow's process view: the real event stream, narrated as business processes a non-engineer can read. The stages light in sequence as it actually happens.

Then the last ask: *can I click a box and see the detail? Can the colour change the moment an event happens?* Yes and yes. Each stage became clickable, opening the specific events behind it; each customer's journey could be followed as an individual thread through the process. The choreography lit as a customer really moved through it — the invisible machine, at last, telling its own story to anyone in the room.



Left: stages lit and cooling as events land. Right: a single customer's journey traced through the process. Click any stage to see the events beneath it.

A system that can narrate itself to a stranger is a system that finally understands what it's for.

The lesson. Observability for engineers is logs; observability for the business is a story. The same event stream, told as processes a layperson can follow, turns an invisible distributed system into something a customer, a buyer, or a regulator can simply watch.

Rules without redeploys, and the audit

In telecom, every new product brings new business rules. The question that decides whether you have a product or a project is: does a new rule mean new code?

— from the build log

The question arrived plainly, the way the best ones do: *"New products have new business rules. Is it new code every time, or is there a rules page?"* The honest answer, at that moment, was both. Prices, bundles, allowances, promo codes — all data, changeable in the console with no redeploy. But a genuinely new *kind* of rule — "max two SIMs per order," "these two products can't be bought together" — was a code change. There was no rules engine, and pretending otherwise would have broken the book's one law.

So the thirty-first component was born: **policy**. Rules authored as pure data — JSON-logic conditions an operator writes in the console — evaluated at order time by a small engine that can never execute code, only answer true or false. The proof was the point: author a "max 2 per order" rule through the API, watch an order for three get refused with the rule's own message, watch an order for two pass, then *disable the rule with a row* and watch the same order for three sail through. The rule appeared, enforced itself, and vanished — and the deployment never moved. The same engine then learned to answer a second question: not just "may they?" but "at what price?" — pricing rules that adjust a subtotal by segment, by cart, by verified identity, stacking like a customer would expect, landing on the bill as labelled adjustment lines. And the bundle model grew the formal cardinality the standard always offered: mandatory components, optional add-ons a customer may include, and "pick exactly N of M" choice groups enforced at order time.

A rule you can add, enforce, and retire without a deployment is the difference between selling a product and selling a promise of future engineering.

And then, because the machinery was all data and seams, came the reckoning the second chapter had promised: *run the official conformance kits against everything else*. The kits were museum pieces — Node-16-era runners that mangled URLs on modern machines — so the first day was spent building a harness that made their numbers trustworthy at all. Then, component by component: measure the real gap, reshape to the spec, redeploy, re-run, and re-run the *application's* own tests so conformance never broke commerce. Shopping cart. Party roles. Stock. Usage. The customer bill — nineteen thousand assertions on that kit alone. Consumption reports. Each driven to zero failures, each committed with its number.

The most valuable finding was the two components that *refused*. Payment would not accept the kit's empty, amount-less payment — because in this BSS, creating a payment is the authorization, idempotent by correlator, and no conformance score was worth re-opening the door to a double charge. Communication would not send a message with no recipient. Eleven kits green, two gaps kept on purpose and written down in plain sight — because a conformance page that admits "here we are stricter than the spec, and here is why" is worth more to a buyer than a wall of unexamined green.

The lesson. Configurability is a spectrum, and the honest move is to say where you are on it — then move the expensive end (new rule types, new prices) into data. And when an external standard disagrees with a guarantee you chose deliberately, the answer isn't to surrender the guarantee; it's to document the disagreement and let the buyer judge. Eleven zeros and two explained exceptions beat thirteen unexamined zeros.

ACT SIX

The Operator's Week

*Companies get a door of their own, the shop
becomes a home, products learn what they are —
and the whole machine starts speaking Norwegian.*

Everyone gets a door

"Do we have full B2B support in our BSS?" The honest answer was no — and honest answers are where features come from.

— from the build log

Every chapter so far had served one kind of customer: a person. But operators sell to companies, and a company is not a big person. It has members whose lines it pays for, an administrator who is not the operator's staff, and an accountant who wants *one* invoice, not forty. The audit took an afternoon and returned a familiar shape: the party model could hold organizations, nothing connected people to them, and nobody but the operator could manage anything.

The foundation went in the way foundations had gone in all along. Organizations gained a parent; individuals gained a membership; a new customer-side role — the business administrator — could act only inside their own organization, a boundary enforced by live lookups in the party service, never trusted from the request. Billing learned consolidation: members' charges regroup under the company, each line still attributed to the person who used it, and if the party service is unreachable the run fails open to per-person bills rather than billing nobody. A sixth surface appeared at `/biz` — the business console — where Bianca of Acme adds people, orders lines for them, and reads the company invoice.

Two truths surfaced in the proving. The first was operational: the consolidated bill came out empty because billing's machine account lacked one read permission — and the fix did nothing until the service restarted, because machine tokens are cached. That lesson ("grant, then restart, then believe") would pay for itself twice more before the week ended. The second was structural: Bianca's people existed as party records but could not sign in, because every channel keys a customer to their identity-provider subject. So invitation became a first-class feature: adding a person with an email mints a real login through the TMF672 seam — hardcoded to the customer role, so the power grants no escalation — and pins the party's id to the new token subject. Erik signs in and lands on *his* line.

Where should Erik land? The storefront was the obvious answer and the wrong one. "I think B2B my-page should be separate," came the verdict, and the channel split in two by role: the administrator sees the org console; a plain member sees a work line, usage, SIM care, and a note that the company pays — nothing to buy, no personal bills. One channel, two faces, the same pattern the role-scoped consoles had established.

Then money. A company that brings forty lines does not pay list price, and the machinery for that already existed: the pricing-rules engine. The billing run began passing the paying organization into the rules context — a field that simply does not exist for consumers, so a corporate deal can never leak to retail — and "Acme pays 20% under list" became a row of data: 25 becoming 20 in the company's price view and on the consolidated

invoice, to the cent. When the raw JSON condition drew the fair complaint that account managers should not write JSON, the console grew two plain-English rule kinds: pick the company from a dropdown, type the percent. The engine speaks JSON; the humans never have to.

THE BILL SPLITS IN TWO

"No personal bills" held for exactly as long as it took to ask the next question: can an employee buy something the company doesn't cover? Netflix on the work phone. A better handset than the policy allows — "many companies pay a fixed amount, say 8,000; above that the employee pays. This should be configurable." The consolidation model had one bucket per company, and real life needed two.

The design fell out of a single observation: *the payer follows the orderer*. When the business administrator places an order, the ordering service stamps it with a payer — the organization — and the products inherit the stamp for life. When a member buys something themselves, there is no stamp, and the charge is theirs. Billing stopped asking "whose company is this?" and started asking "who agreed to pay for this?", regrouping every product by its payer: company-ordered lines consolidate exactly as before, self-bought extras land on a personal bill, and usage is rated once, on whichever bill carries the person's plan. A one-time backfill stamped every pre-existing company product, so the day the feature shipped, no invoice moved by a cent.

The device allowance became the customer's own dial, not the operator's. The business console grew a "Company policy" card where the admin sets the cap — and the guard around it is the same lesson as every boundary in this book, enforced server-side: an admin may tune *their own* company's policy and nothing else; renaming the company through that door returns a refusal, and a plain member touching any organization sees a 404. Come the billing run, the company invoice carries "Samsung Galaxy S26 (*company share*)" capped at the allowance, and the employee's personal bill carries the remainder as its own labelled line — *above company allowance* — to the cent. The proof is the twenty-second suite: an admin sets the policy in a real browser, a member buys Netflix with her own login, and two bills come out split exactly right.

The week's hard lesson arrived sideways, as they do. Mid-verification, suites that had nothing to do with billing started failing: personas whose pages had worked for days could not read their own records. The trail led back to the identity provider — a realm re-import during earlier work had quietly re-minted user ids and dropped every role granted live, so logins no longer matched the party records that owned the lines, the SIMs, the bills. The fix was to make identity as declarative as everything else: persona ids pinned in the realm files, the business-admin role a composite in the files rather than a live grant, and the drifted users restored to their original ids so years of accumulated state reattached. "Grant, then restart, then believe" gained a corollary: if it only exists live, it is already lost — the files are the system.

The shop becomes a home

A customer doesn't own a list of rows. They own a plan, some lines, a SIM in a phone, and a bill they'd rather not think about.

— design note, My page

The critique arrived as five observations in one message: the page mixed the commercial with the technical, the plan-change menu offered a Samsung as if it were a tariff, bundles hid their contents, and where were the ten- and fifty-gigabyte plans, the one-time data top-ups? Behind all five sat one missing idea: the catalog had no notion of what an offering is. So it learned. Categories — mobile plans, broadband, devices, TV, top-ups — became persisted, first-class facts, and everything else followed from them.

The page rebuilt itself around the MyJio idea: it recomposes around what the customer holds. A card per line of business — the bundle with its components nested, every numbered line with its own SIM, the latest bill with a pay link, a discovery card fed by the recommendation engine. Plan change became a real TMF622 [modify](#) order: same service, same number, validated hard (your own product, an existing plan, like-for-like by category) and blocked while a commitment window is open. A bundle's broadband tier upgrades in place because bundles decompose into per-component products, while the commitment binds the bundle itself. The SIM became real — ICCID and PUK minted at activation, PUK revealed on request, PIN pushed over the air through a pluggable SIM-platform seam — and a five-gigabyte top-up became a purchase that genuinely grows this month's meter, delivered through the event stream and idempotent per order.

Complicated products got their showpiece: Family Max. Two fixed components, "family lines — pick one or two" of three plans, a configurable phone whose colour and storage ride the order, up to two streaming extras, a twelve-month term, a base price plus a one-time fee while every chosen option prices itself. The storefront configurator learned pick-N-of-M — checkboxes that cap themselves live, an add-to-cart that gates below the minimum — and the order API enforces the same limits, because the UI is convenience and the order is law. And when the fair question came — "does a product owner need to use JSON?" — the console grew a bundle composer: components marked included or optional, choice groups with their pick-ranges and defaults, terms and categories as plain fields, and a promise that anything it doesn't understand survives a save untouched.

Then the products that aren't lines at all. Netflix on the operator's bill is sold like any offering but fulfilled with the *partner*: the orchestrator now reads the catalog category to decide fulfilment — partner services mint an entitlement code through a partner seam, security products activate as resource-less features, insurance provisions nothing and simply bills. The proof was cinematic in the most literal way: reviewing stills of the demo film caught the top-up minting a phantom phone line with its own SIM, and the film's Norwegian take caught the partner machinery failing for the second tenant — one missing per-realm grant. Both fixed, both reproven, both on camera in the next cut.

THE SHOP GETS EYES

A phone sold as a row of text is a spec sheet, not a phone. The question behind the fix was architectural before it was visual: where should product imagery *live*? Retail has a whole product category for this — the PIM, product information management — and the tempting answers were both wrong. Building a full PIM into the catalog would weld a content factory onto the commerce spine; requiring one would make a heavyweight system a prerequisite for showing a picture of a Samsung.

The answer was the same shape as the PSP, BankID and the porting clearinghouse: a seam. The catalog stays the only face a channel ever sees — imagery is the standard TMF620 attachment list, facts are spec characteristics — and where the content *comes from* is per-tenant configuration. By default it's the document component the stack already had: the product owner uploads gallery shots and colour variants on the offering form, no JSON, no extra system, and that is genuinely enough for a telco's device range. The names carry the semantics: a `gallery-*` attachment builds the product page's photo strip and puts phone *pictures* in the bundle configurator; a `variant-Icy Blue` attachment makes the hero photo follow the colour picker; a characteristic marked non-configurable renders as an "About this device" table — display, camera, battery — instead of a dropdown.

And when an operator already owns a PIM — years of curated content, a merchandising team that lives in it — they keep it. One registry entry points the tenant's catalog at their system, and offering reads resolve imagery live, keyed by product *name*, because their PIM has never heard of our UUIDs. Cached for a minute, failing open to whatever the catalog stores: content must never block commerce. Nova became the living proof — a "Nordic Phone X" created in its catalog with zero artwork arrives in the shop with a full gallery served from Nova's own system, through the same gateway, on the same storefront build. The twenty-third suite walks both paths and can't tell them apart from the browser, which is precisely the claim: no PIM required, any PIM welcome.

The pictures raised the next question within the hour: "sometimes prices vary based on phone colour as well — is that going to create many SKUs?" Retail's answer would: one catalog row per colour per storage per device, each with its own prices, stock and campaign references, multiplying until nobody dares touch the range. The TMF answer is a price component with a condition — `prodSpecCharValueUse`, "+2.00 a month when colour = Titanium Edition" — on the *same* offering, the same spec. The storefront learned to make the price follow the pick exactly as the photo already did: choose Titanium and the hero swaps *and* the premium row appears; choose Phantom Black and both revert. The billing run learned the same trick from the other end, rating each customer's product against its *configured* characteristics — two neighbours buy the same Samsung, one bill says 39.49 and the other 37.49, and no SKU was ever minted. And because the configured picks now ride the pricing context as plain values, marketing can run a campaign on a colour: one rule — the console has a preset for it — puts "Titanium launch: 10% off" on the cart preview and the invoice of Titanium buyers only. The twenty-fourth suite proves the whole chain, catalog to invoice, to the cent.

THE CATALOG LEARNS TO LISTEN

The composer had removed the JSON; the next request removed the forms. "I want a feature like a copilot — the product owner chats about the type of product they want to create, it tells them how, and actually creates it if they say yes." The temptation with such a feature is to hand the model the keys. We did the opposite, and the architecture is the story: the model only ever *proposes*.

The conversation runs through the intelligence component like every AI feature before it — the same any-model seam, the same audit ledger, and the same privilege stance as the CSR copilot: the console sends the catalog context the owner's token can already see, and the endpoint fetches nothing on its own authority. The owner types "I want to sell a streaming service" and the copilot answers from the modeling rules this book has been accumulating for six acts: that's a *partner service* — activation will mint an entitlement code with the partner, billing carries the charge — and what should it cost? The reply comes back as a structured proposal: a spec, a price, an offering with its category, rendered not as JSON but as a card — *the copilot will create: spec · StreamPlus Service; price · 9.99 EUR/month; offering · StreamPlus [Partner services]* — with one button.

"Yes — create it" is where the design earns its keep. A validator checks the proposal the way the composer would (duplicate names, dangling references, prices that aren't positive), and then a deterministic executor applies it in dependency order — spec, prices, offerings, bundle links — using the product owner's *own* token, rolling back in reverse if anything fails halfway. The model never holds credentials; role-based access, tenant isolation and the audit trail all hold because the write path is ordinary code making ordinary TMF620 calls. The twenty-fifth suite closes the loop with the only proof that matters: a customer *orders* the copilot-made product, and the orchestrator mints a real partner entitlement code — the copilot's category choice drove genuine fulfilment. Its second scenario chats a four-part kids-smartwatch bundle into existence — device, child plan, optional GPS tracking, a twelve-month term — applied in one confirmation. And because the deterministic stub provider carries both scenarios, the whole conversation runs in CI with no model anywhere: chat to catalog, keyless, verified.

Kroner and tongues

"In Norway it will require Norwegian, and NOK as currency."

— the requirement, in full

The audit was two sentences long. Currency: eighty percent there, because every amount in the TMF models had carried its unit from the first chapter — EUR appeared in code only as a fallback. Language: zero percent there. And so the fix followed the house's oldest rule: the tenant manifest, which already told every channel its brand name and colour, learned two more words — `locale` and `currency`.

The channels grew a translation layer with one deliberate trick: English strings *are* the keys, so an untranslated string falls through as itself and the English tenant risks nothing. Money formats through the platform's own internationalisation from each price's own unit — "299,00 kr/md." in Norway, the historical English format untouched because a dozen test suites assert it. Keycloak's built-in Norwegian switched on per realm. And Nova Telecom, the white-label tenant that had existed to prove isolation, became a real Norwegian operator: NOK prices, a tariff ladder, a top-up called Datapåfyll, and personas to walk it — Nils on Min side changing plans in Norwegian, keeping his number at 199 kroner.

B2B demanded the same parity and delivered the best find of the week: the business console had its identity provider hardcoded to the first tenant's realm. It could never have served a second operator — not mistranslated, but structurally monolingual. Wired into the manifest like every other channel, it mounted at Nova's own hostname; Birgit Berg of Fjellheim AS runs her people and lines in Norwegian and reads one consolidated invoice at 299 kroner. When the header still read "BEDRIFT," the correction came from the person who knows: Norwegian operators call it *Min bedrift*. Renamed. The whole B2B machine — boundaries, consolidation, invitations — worked for the second tenant with zero backend changes, which is the sentence the first chapter was always trying to earn.

The tenant-isolation suite now walks an entire purchase in Norwegian — Legg i handlekurven, Til kassen — and the demo film ends its Norwegian act with a plan change and a company invoice, in kroner, before Live Flow rolls the credits. One build. Every language. Every currency. The manifest did it, not a fork.

Proof against tomorrow

"Let me know if our BSS is PQC proof." Five words, and suddenly the audit was of mathematics itself.

— from the build log

Every audit so far had asked whether the system did what it claimed. This one asked whether its arithmetic would survive an adversary that does not exist yet. Post-quantum cryptography sounds like a problem for cryptographers, but for a BSS it is a problem of *inventory*: where, exactly, does cryptography live in this thing? The grep took minutes and returned the most reassuring answer available: almost nowhere. A secure random generator for PUKs and invitation passwords. A SHA-256 challenge in the login dance — hash-based, and hashes only lose half their margin to a quantum computer. And that was the entire list, because identity, transport and payment crypto had all been pushed behind seams in chapter one, back when "vendor-neutral" was a conviction and not yet a security posture.

The vulnerable pieces were exactly two, and neither was ours. The identity provider signs every access token with RSA — the algorithm Shor's algorithm eats first. And transport is classical TLS, which matters more than it sounds: an adversary can record encrypted traffic *today* and decrypt it the day their quantum machine boots. "Harvest now, decrypt later" is the one post-quantum risk with a present-tense clock, and it is entirely a deployment concern.

So the honest answer went back in two sentences — not PQC-proof, PQC-*ready* — and the follow-up question wrote the afternoon's work: so what's ours to do? Three things, and two of them shipped the same day. The Helm chart grew an edge: TLS 1.3 only, with controller values that switch the key exchange to a hybrid — a classical curve and ML-KEM-768 riding the same handshake, so clients that speak post-quantum get it and everyone else falls back gracefully. And the SIM PUKs, sitting in the database as an honest dev shortcut, got a vault: AES-256-GCM with the card's own ICCID bound into the ciphertext, so a stolen row cannot even be replayed onto another card. Old plaintext rows quietly upgrade themselves the next time their owner asks to see them. AES-256 is already comfortable against Grover; the fix was classical hygiene wearing a quantum justification.

The third item was the fleet itself: thirty services on a JVM that had never heard of ML-KEM. "So what's stopping the backlog?" Nothing, it turned out, except doing it — and doing it honestly meant meeting three of Spring Boot 3.5's edges in one afternoon: Flyway had modularised its databases, so twenty-six services booted to "Unsupported Database: PostgreSQL" until each carried the new module; the gateway's Spring Cloud train refused the new Boot on principle until it was bumped two generations; and one test dependency fell out of the managed BOM. The subtlest lesson wasn't in any pom: a single transient image failure aborted the parallel build and left a *mixed* fleet that still reported healthy — old containers serving happily under new build logs. Trust `java -version` inside the running container; never trust the tail of a build.

Then the definition of done: all thirty services healthy on Java 25 LTS, and the full twenty-one-suite sweep — every channel, every tenant, every language — green. The remaining migration, the IdP's RSA signatures becoming ML-DSA, waits on the OIDC ecosystem, and when it lands the services will not notice: they have never hardcoded an algorithm, they follow the issuer's published keys. That is the whole argument, and it doubles as the marketing copy precisely because every claim in it has a receipt: a BSS that embedded its own crypto would face a post-quantum rewrite. This one faced a checklist, and the boxes that were ours to tick are ticked.

C O D A

What it means

*One person, one AI, one discipline — and a system
that used to take a team.*

The family ledger

"I think before developing you should think about plan, do some research on ux and flow by other big companies like Orange, Jio and the Nordic CSPs... because I really want to give a good user experience for end customers."

— the operator, pulling the handbrake

The household had started as billing plumbing: a person could pay for another person, consent-gated, one family bill with per-person attribution. Correct, tested, honest — and not an experience. The operator's own click-through proved it. Two subscriptions on a page with no hint of whose they were; a link that promised "open their page" and opened the wrong one, because a login redirect quietly dropped the deep link on the floor. The bug was one line. The verdict was bigger: stop patching, go read how the operators people actually live with present a family.

So the research came first, for once, and it converged fast. Jio's parent SIM is a hub — the parent's app lists every linked number and drills into each one, and adding a member sends the OTP to *their* phone, consent built into the gesture. Orange lets a second parent co-pilot the children's accounts. Verizon has the cleanest role grammar in the industry: one Owner, up to three Account Managers who can do nearly everything except manage each other, Members who see only their own line. T-Mobile's Family Allowances let the owner set caps the managed lines can see but not change. And Google's Family Link solved the purchase question years ago: the child taps "buy", the parent gets the ask, nobody hits a dead end. None of them scatter labels across the member's own page. All of them build a family *place*.

The family hub followed in three phases, each shipped against a live suite. Phase one was the place itself: a Family tab where the payer — the *owner*, in Verizon's grammar — sees every member as a card with a role chip and the services the family funds, and where "my wife can also be the admin" became a button. Promotion is the owner's alone; an admin sees the whole family and orders for any member, but the payer stamp on that order is the *household payer's*, never the admin's — admin is authority, not a wallet. The suite proves the sentence: the wife orders the son's plan from her own hub, and the bill lands on the owner. Privacy stayed proportional the whole way down. Through the family window an adult shows only what the family funds; the son's own Netflix, bought with his own money, never crosses the glass. Paying is not watching.

Phase two took Family Link's insight and T-Mobile's caps and fused them onto the order pipeline. Each member can carry a monthly top-up allowance in euros, set from their card on the hub. Inside it, a child's top-up completes instantly on the family bill and the payer just hears money moved. Above it, the order parks in a state the fulfilment layer ignores by its own guard — *held* — and an approvals inbox grows at the top of the hub. Approve, and the state flips; the machinery picks it up exactly like a fresh order. Decline, and the requester gets

a kind sentence in their inbox instead of a dead end. Adults are never blocked: over the cap they simply pay from their own pocket, as any adult may.

The Nordic move

Then came gifting, and with it the question that mattered more than the feature. In Norway, a CSP rolls unused data forward automatically and lets a customer gift it to *anyone on the network*. "Are these features configurable," the operator asked, "or does everything need to be recoded?" The honest answer was: hardcoded, today. The right answer took an afternoon: gifting scope, share caps, giftability, rollover eligibility and rollover caps all moved onto the plan's spec characteristics — product data, read through a thirty-second cache — and the policy engine grew a third rules-as-data family, *gifting* vetoes, for whatever the levers can't say. The proof was theatrical on purpose: the test itself, acting as Nova's staff, put `giftScope=tenant` on Nova's plan spec — a catalog edit — and two complete strangers gifted gigabytes to each other on one tenant while the other tenant, on the same running binary, still answered "no household link to that person." The feature nobody hardcoded, running both ways at once.

Gifting by phone number closed the loop, because that is how people actually know each other. The number pool had held the answer all along — every assigned MSISDN already mapped to its holder in the fulfilment layer — so a typed number resolves server-side, agents and machines only, and the confirmation echoes the number back, never the stranger's name. Names cross by consent; numbers are how you address a gift. Phase three made age real data instead of an assumption — a birth date on the party record, an adult joining as a member where an undated dependent stays safely a child — and gave the guardian's window its meters, for children only. The child's card on the hub carries a live usage bar. The adult's never does, even when the family pays. The suite asserts both sentences, because the second one is the product.

The library

"I also need a knowledge management feature... so customers or CSR or sales can query if they are wondering on stuff within BSS or channels. How big is that to build?"

— the operator

By now the platform had thirty-one components and a family saga, and the honest gap was that nobody could *ask it anything*. The answer became component thirty-two, and it followed every rule the previous thirty-one had taught. Articles and FAQs as data, authored in the console, tenant-walled by row-level security, searched by the database's own full-text engine — no new infrastructure, because Postgres was already load-bearing. One library, with one law: *who you are decides what you see*. A customer searching "gift" gets the FAQ shelf on the shop's Support page — answers first, tickets second. An agent running the same search gets the cheat-sheets on top: the household role grammar, what a held order is and who may release it. A product owner opens the console and finds the shelf this book's machinery deserves: how to create a product — spec, offering, price, in that order — how bundles compose, which spec characteristics flip gifting network-wide, how to talk to the copilot. The how-to-build-products library lives where the products are built.

The layer that made it more than a search box reused the oldest seam in the system. Ask a question in plain words at the CSR desk and the intelligence component retrieves first, then lets the model answer *from the retrieved articles only*, sources named, honest "I don't know" when the library is silent. The retrieval runs with the asker's own token — not a machine identity — so a customer's question can only ever be answered from articles a customer could read. The audience filter isn't advice to the model; it is the shape of what the model is allowed to see. On the keyless stub the answer points at the right article; with a real key it becomes synthesis. Either way it cites its shelf.

The week's small confession belongs in this chapter, because it is the same lesson at desk height: the CSR customer search had been an *exact family-name match* since the day it was born. Typing "gina" — a given name, lowercase — found nobody, in a system that could split a bill three ways. The fix was ordinary (any case-blind fragment of a name, an email, an address line; results settling three hundred milliseconds after the last keystroke) plus one flourish the number pool made free: paste a phone number into the same box and the holder's 360 opens, their numbers now sitting in the header where a ringing phone needs them. Features are load-bearing only if the person at the desk can find the customer they're talking to.

The shop that pays attention

"Is building a CDP for this BSS as a BSS component wrong from an ODA perspective, or otherwise?"

— the operator, asking the right question before the first line was written

Personalization arrived the way every good feature had: as a question with a trap in it. Connect Google Analytics and social signals for visitors, fuse them with BSS data for customers, personalize the pages and the offers — every operator wants it, and the obvious way to build it is wrong. The obvious way is to become a CDP: an identity graph, an event lake, audience machinery. The research said the field always builds the same pipeline — events, identity, profile, segments, a decision, an experience — and the architecture said the pipeline does not have to live in one box. So the answer to the operator's question became the design: a BSS should be *CDP-ready*, never a CDP. The enterprise identity graph belongs behind a seam, like the charging system and the PIM before it. What the BSS legitimately owns is small and load-bearing: the first-party, service-context profile — and the consent that gates it.

Consent went in first, as structure rather than ceremony. A loud, brand-framed card on the shop's front page asks one question in plain words, decline exactly as prominent as accept, and the promise underneath is mechanical: nothing is stored before the answer, a rejecting visitor generates *no rows at all* — the suite fires a rude, consent-ignoring event straight at the API and asserts the profile holds nothing — and revoking consent deletes what was held. "Privacy choices" reopens the card forever, because changing your mind must be as easy as consenting was. The banner's promise is real, and tested, which is the only kind of privacy copy worth printing.

Then the loop: a consenting guest browses a Samsung, and the home page says so — "because you were looking at Devices" — with the Devices rail leading. The decisioning went where all behaviour goes in this system: the policy engine grew its fourth rules-as-data family, *personalization* experiences, so an operator writes "visitors interested in Devices see this banner and that pinned offer" in the console and deletes it just as easily. Signing in performs the one CDP trick worth stealing — the stitch: this browser's guest interests now belong to this customer, by consent, and the TMF680 recommendation engine folds them into its ranking. The suite's customer browsed one phone as a stranger and got Devices at the top of "Recommended for you" as a customer. BSS truth plus browsing, one answer.

The seam flows both ways

The analytics seam kept the promise the OCS and PIM chapters made: bring your own. Nova configured `ANALYTICS_PROVIDER=ga4` and its consented events flowed to "its own GA4 property" over Google's real Measurement Protocol — a mock playing GA4 in dev, exactly as mock-ocs plays the charging system — while genalpha, on the same running binary, stayed first-party only; the test asserts a genalpha visitor never appears

in nova's property. And the seam pulls as well as pushes: audiences the platform computes come back as rule-addressable segments, campaign tags classify the channel (instagram is social, google is search), and a console preset turns "an analytics segment sees this experience" into one form. On top of it all, next best offer: the CSR presses one button, the TMF680 ranking supplies the candidates, and the model's only job is the WHY — "they have been browsing Devices; this is the strongest match" — grounded, sourced, never inventing an offer that is not on the shelf. The console got an Insight ledger so the operator can SEE what the platform holds and under which consent — read-only by design, because the visitor owns the data; the operator only gets to look.

The growth engine, hands-on

"Why is a BSS only product and billing? Most of the data lives here. I want to challenge the status quo."

— the operator, and then the operator's whole roadmap

This chapter is the manual the thesis earned: how a product owner actually RUNS the marketing engine, end to end, with nothing but rows in the console. Everything below is data — no deploys, no vendor onboarding, no CSV exports — and every step has a suite behind it. It reads as one continuous play because that is how the engine was built to be used: the shop pays attention, a segment forms, a campaign speaks, the readout tells the truth, and every word said lands on the customer's record.

1 • A segment is a noticed interest

There is no segment builder, because there is nothing to build. A consented visitor who keeps returning to the *Devices* category IS the Devices segment; the login stitch makes the membership customer-grade. In the console's **Insight** tab the operator can see the ledger; in **Audiences**, the names the tenant's own analytics computed (imported live over the GA4 Data API's `runReport` wire) sit beside the first-party interests, equally addressable. Any of those names can go straight into the next step.

2 • A campaign that can prove itself

Console → **Campaigns** → New: name it, put `Devices` in the segment field, write the message ("`{code}` inserts the promo"), and — this is the part the industry skips — set **Holdout 10%** and a seven-day conversion window. Execute. A deterministic tenth of the segment is ledgered and deliberately NOT messaged; they are the control group. A week later the detail row does not say "reached 4,000" and stop; it says *treated converted at 22%, holdout at 8%, fourteen points of lift, worth 6.10 per customer/month* — the conversion carries the catalog's own monthly price of what was actually ordered, so lift reads in currency. Under five holdouts, the row says "an anecdote, not a measurement," because a readout that flatters is worse than none.

3 • Two pitches, one honest verdict

Add **A/B arms** as a JSON field — two to four variants of `{name, subject, content}` — and the treated group splits evenly and deterministically. The stats read per arm, and the verdict is a two-proportion test that refuses flattery: under thirty people per arm it answers "the split is an anecdote — keep the test running," and it crowns a winner only at 95% confidence. The suite's favourite assertion: six customers, one conversion, and the engine declining to declare victory.

4 • Journeys that read before they speak

Journeys are sequences as data: `[message, wait 3 days, branch, message]`. The journey-level conversion event is also the ALWAYS-ON exit rule — the customer who buys during the wait leaves from any step and never gets "10% off!" the day after paying full price. The `branch` step asks insight a question at send time — *is this customer in the Devices segment?* — and speaks the matching pitch: the gadget lover hears the accessory line, everyone else the generic one, nobody hears both. An unreachable insight fails SAFE to the generic side, on the principle that the right-but-late message beats the wrong one.

5 • Guardrails are settings, not etiquette

The **Guardrails** tab holds two numbers and a window: *max marketing messages per customer per N days* (campaigns skip the capped, journeys postpone an hour) and *quiet hours* in tenant-local time (campaigns reach nobody inside the window — and the same execute reaches them in the morning, skipped never lost; journeys park to the window's end). Braze's own retrospective said teams forget optional guardrails, so here they are model, not manners.

6 • The message leaves the building — and answers back

With `DELIVERY_PROVIDER=esp` a tenant's messages ALSO ride its own email provider — SendGrid's v3 wire shape, address looked up live from party management, the in-app inbox always the source of truth. The provider's receipts come home over its webhook: delivered stamps the message, a bounce puts the address on the suppression list and the engine never emails that door again — in-app continues, compliance is not optional. Segments flow outward too: one call pushes a segment to the tenant's social platform as SHA-256 hashes (Meta's Custom Audience schema — not one address in the clear), and the platform's lead-gen forms flow back through the same seam into TMF699 salesLeads, idempotent on the platform's own id. Social ad → lead → qualify → opportunity → quote → order, one system.

7 • Every word lands on the record

The oldest component in the sales story turned out to be the keystone. TMF683, the interaction log built back in the CSR chapter, became the OMNICHANNEL record: every message the communication component sends — a blast, a journey step, an order notification — mints a touchpoint on the customer's timeline automatically, tagged with its source. And because the TMF683 door is a plain open POST, anything else in the operator's landscape — the legacy CRM, the retail till, the field-tech app — logs its calls and visits into the same timeline. The CSR opens the 360 and reads, in one place: *campaign pitch Tuesday (outbound, in-app, via communication), order confirmation Wednesday, a phone call Thursday (inbound, via legacy-crm, "prefers email, has 12 vans")* — and speaks like someone who was paying attention. That timeline is what "omnichannel" means when it is real: not a slogan about channels, but one memory shared by all of them — the shop, the CSR, the campaign engine and the systems the operator already owns, none of them asking the customer to repeat themselves.

A marketing tool proves it moved money. A sales tool remembers every conversation. A BSS that does both, natively, is no longer missing the point of its own data.

The load-bearing wall

"Make it work, make it right, make it fast."

— Kent Beck

Thirty-three components. Six customer surfaces, one of them wearing two faces by role. Forty end-to-end suites, all green in a real browser. Eleven official conformance kits passing with zero failures — and two components that fail their kits on purpose, with the reasons written down. Business rules, prices and corporate deals that change with a row, never a redeploy. Isolation proven against a real database — and then proven again in Norwegian, in kroner, B2C and B2B, from one build. Plans that change without losing the number; SIMs with a PUK you can ask for; top-ups that genuinely grow the meter; Netflix fulfilled as a partner entitlement because the catalog category decides what a product is. A bundle composer so product owners never touch JSON, and an order API that enforces what the composer promised. A national number-portability flow, both directions. A quantum-safe edge and encrypted card secrets, with the one remaining post-quantum migration parked at a seam. A network that heals itself and leaves a receipt. And a film — re-shot after every feature — that keeps catching real bugs, because everything on camera is real.

Look back and the three convictions turn out to have been one conviction wearing three coats. **Vendor-neutrality** — any IdP, any database, any broker, any model, any processor, any porting authority — was possible only because every dependency was a seam, and every seam was verified with a real implementation behind it. **Honesty about mocks** was survivable only because the mocks were named and the seams were real, so the day a real Stripe or a real NRDB plugs in, nothing above it changes. And **verification** was the wall that held both up.

That wall is also the answer to the question this book is really about. It was built by one person and an AI — Claude Fable 5 — at a pace no solo human could sustain and no unverified pair could survive. The model's breadth was extraordinary; it held the whole system in view and wired thirty-one components without losing the thread. But breadth without an arbiter is just confident wrongness at scale. The tests were the arbiter. Not the model's fluency, not my intuition — a browser, a real database, an official kit, agreeing or refusing.

Give the best model the whole system to hold, and a discipline that catches it when it's wrong, and one person can build what used to take a team — if they never stop verifying.

That is the memoir's whole argument, and the reason its title is an instruction rather than a boast. The future of building software isn't "the AI writes the code." It's a human holding intent and judgment, a capable model

holding execution and breadth, and a ruthless verification discipline making the partnership trustworthy. A thirty-one-component, multi-tenant, standards-conformant BSS, built this way, is the existence proof. Every claim in these pages has a commit, a test, or a screenshot behind it.

Verify everything.

What happens when the customer calls

"What if the customer wants to..."

— the operator, eleven times in a row, and every time the honest first answer was "they can't, yet"

There is a genre of question that sorts real BSSs from demos, and the operator asked it eleven different ways in one relentless stretch. What if the customer forgets their PIN? Loses the SIM? Wants a new number? Wants to pause the subscription for a month abroad? Can't pay the whole bill? Wants to upgrade mid-cycle — who pays for the half-used month? Wants a refund? Wants to leave and take the number? Wants to hand the work line to a colleague? Says the internet is slow? Wants their daughter on the plan? Wants the bill on payday instead of the first? None of these are features a launch demo shows. All of them are Tuesday at a call center.

The gauntlet

Each question became the same loop: research how real operators handle it, build the lifecycle act, notify the customer, land it on the timeline, prove it in a suite. A changed PIN is never silent — the owner hears about it, because a silent credential change is how fraud reads. A lost SIM quarantines the old card before minting the new one. A paused subscription pauses charging at the OCS and comes back by itself on the promised date. A number change quarantines the old number with its story written down. And the security posture inverted properly along the way: self-care endpoints stopped demanding staff powers and started demanding OWNERSHIP — any authenticated customer may act on their own line, and a foreign line answers 404, never 403, because even "forbidden" leaks that there is something to forbid.

Money got honest

The billing half of the gauntlet turned the engine segment-based. A mid-month start pays from its first day, labelled "(from the 16th, 15 days)". A mid-cycle plan change bills each plan for exactly its own days. A cease bills to the end date and no further. Underneath them all, one invariant the suites enforce from every direction: *no day is ever billed twice* — every billing run clamps its period to the day after the last covered one. On top of the honest bill came the honest hardship path: an unpaid bill splits into installments that sum exactly; a missed one earns one reminder, then a grace, then the broken plan lands on a dunning worklist with the remainder due through one visible door. A disputed charge pauses collection while a human decides, and the decision — credit or uphold — is written to the customer in either case. Payday alignment closed the arc: pick a day, the cycle

follows it, and the bridging stub bill covers exactly the gap days, because the no-double-billing invariant holds even across a change of rhythm.

*A demo sells the happy path. An operator lives in the other eleven.
The system is only real when "what if the customer wants to—" stops being a stump and starts being a suite number.*

The bill leaves the building, and the money comes home

"I see that bills don't open as PDF... and the format we support in Norway is EHF, whereas I guess it's EDI in Europe. Check and make this a configurable option."

— the operator, half right, which is the useful kind of wrong

The research corrected the premise first, which is what research is for. Europe's e-invoicing standard is not EDI: it is EN 16931, one semantic model with two syntaxes — UBL 2.1 and UN/CEFACT CII — and Norway's EHF is a national profile of Peppol BIS Billing 3.0 riding the UBL syntax, with one national demand that matters: a payment reference, the KID, stamped into every invoice. EDIFACT, the operator's guess, is the pre-XML legacy a trading partner might still demand — real, but a follow-up, not the standard. That correction shaped everything built next.

Formats became rows

A format profile is policy, not logic: which syntax, which CustomizationID, whether a payment reference is required. So profiles became database rows a tenant admin edits in the console — and the suite proved the rows are load-bearing by editing one and watching the very next e-invoice change on the wire. Australia arrived as a seeded row. Germany's XRechnung was typed into a form. EDIFACT and Factur-X — the hybrid PDF with the machine-readable CII embedded inside it — arrived as two more syntaxes behind the same rows, proven by flipping a row and flipping it back. Adding a country stopped being a release and became an INSERT.

Delivery grew a ledger, and the buyer answered

Fire-and-forget was not how money documents travel, so cutting a bill began writing a delivery ledger row in the same transaction — the outbox property: if the bill exists, the intent to deliver it exists. A relay drains the ledger with backoff; the suite knocked the access point over twice and watched the third attempt land, exactly once. Then the buyer's answer joined the same row: a Peppol Invoice Response — accepted, rejected with the buyer's words, paid — so one row tells the document's whole life. The buyer's "paid" is a claim, though. Money is only real when the bank says so.

And the money came home

The bank says so in dialects. ISO 20022 camt.054 XML. Norway's Nets OCR giro file — eighty-character fixed-width lines, amounts in øre, the KID right-adjusted in columns fifty-one to seventy-five. America's BAI2

lockbox. All three land on one webhook that detects the dialect from the first bytes, reduces it to credit entries, and matches each entry to a bill by the payment reference the invoice carried out. A clean match books a real payment and settles the bill through the same guarantee path a card uses — and the customer is thanked, because a bill that closes itself silently reads as a glitch. Everything unclear — unknown reference, wrong amount, a re-sent file — parks on the unapplied-cash worklist with its reason. Money is fail-closed: never guessed, never dropped.

The bill is a conversation in three directions: the document out, the buyer's answer back, the money home. A BSS that only prints the number has done a third of the job.

Selling through someone else's store

"In Norway, chains like Elkjøp and Power sell the CSP's products."

— the operator, describing the entire physical distribution of a real MVNO in one sentence

A CSP without stores of its own makes the electronics chains its shelf. That model has three shapes and the system grew all of them. The DEALER CONSOLE, a sixth channel, where a clerk whose organization holds a dealer agreement sells at the counter. The STARTER KIT, a SIM in a box with an activation code — and, decisively, the dealer attribution baked into the kit itself, so a box sold like a chocolate bar still credits the store when the customer self-activates on their sofa, no partner system ever involved. And the PARTNER API, because Elkjøp and Power run their own tills: the agreement row names an OAuth2 client, the credential IS the dealer, and a rate limiter at the edge keeps one chain's retry storm from ever crowding out another.

Being a dealer is a row, not a role. A clerk's power is their org membership, checked live; a foreign dealer sees 404-shaped nothing, because attribution is also isolation. And the chains sell their OWN phones — their stock, their fiscal registers, their kassasystemlova obligations — so the device rides the sale as context on the commission entry, never as a billable item. The BSS is many things; a certified cash register is deliberately not one of them.

Commission moves like money

The money-out ledger inherited the money-in discipline. Commission accrues PENDING at activation, because Norway gives every customer fourteen days to change their mind. It hardens to EARNED when the withdrawal window closes. And when a customer uses their angrerett on day three, the pending entry claws back with the reason written on it — not in an email thread, on the row. The suite runs the whole arc on dev clocks: sold, pending, earned, and one honest clawback.

A dealer channel is attribution you can audit, money that can change its mind, and walls between chains. Everything else is a login page.

The third operator in an afternoon

"Can our BSS host an MVNO?"

— the operator, asking the question the architecture had been quietly answering all along

Two operators on one deployment had been the proof of multitenancy. The MVNO question exposed the asterisk: both were baked into every service's configuration as build artifacts. A third would have meant editing thirty-one files and rebuilding the world. So the wall came down in three moves. The tenant fleet became a FILE — one shared registry, the union of every service's per-tenant fields, mounted everywhere; adding an operator became appending a block. Then a script did the afternoon: clone the template realm, append the block, restart, seed — and suite forty-nine timed it at a hundred and twenty-five seconds from nothing to a customer whose first bill was already correctly PRORATED in her own currency, because a newborn operator inherits honest billing for free.

Then the restart went away

Every service grew a small refresher that re-reads the shared file on a timer, and the security resolvers — which had always built their issuer trust lazily — honored a newborn tenant's first token with no further work. That made the operator a FORM: five fields in the admin console — id, brand, locale, currency, color — and the running fleet grows a tenant. Eleven seconds from a blank form to a served customer, zero restarts, zero builds. The white-label came free, because the gateway had always served each storefront's manifest from the registry by hostname: the moment the refresher ticked, `shop.aurora.localhost` wore Aurora's name and color.

The boundary, written down — and the day it bit

Live mutation followed: change a serving operator's brand or bank token in the file and the fleet follows within one interval. But identity — the id, the issuer, the key endpoints — stays restart-territory, enforced in code, because the security caches key by issuer and a boundary you can name beats one you discover. The guard on the mutation endpoint exists because its absence bit within the hour: a failed test run renamed Nova "Hijacked" through the unguarded door, and the refreshers loyally carried the vandalism fleet-wide — then healed it just as fast when the file was corrected. Live mutation cuts both ways. That is why seed operators are protected, form-born operators are form-mutable, and the incident is in this book instead of buried.

Hosting is the multitenancy thesis with a deadline. Two operators prove the walls; the third one, born in an afternoon and rebranded while serving, proves the business.

The channel the law designed

"The call centers who sell subscriptions are generally outsourced partners with their own dialling system... is our BSS catering to them too?"

— the operator, and the answer was "structurally, yes — legally, not yet"

Outbound telesales was the channel where the law wrote the architecture. A call center is, structurally, a dealer with a phone — the same agreement row, the same org-membership agents, the same machine credential for their dialer, the same commission ledger — and all of that carried over untouched. What telesales ADDED was three things the dealer channel never needed, and each one is a legal requirement wearing an API.

The wash

Norway keeps a reservation register: a citizen who says "do not call me" may not be called, and certainly not sold to. So the wash runs FAIL-CLOSED before any offer exists — a per-tenant seam to the register, and three refusals that each say why. A reserved number refuses with the reason. An UNREACHABLE register refuses too, because "we couldn't check" is not consent. And a tenant that never configured a register has no outbound channel at all: the operator opts IN by configuring the wash, never out. The dial list inherits the same honesty upstream — the dialer pulls its audience from the same consented segments campaigns use, every number washed at export, and the reserved citizen is EXCLUDED AND COUNTED ("1 reserved excluded", said out loud on the agent's screen), never listed.

The offer that is not an order

Under angrerettloven, a consumer agreement made on an outbound call binds only when the customer confirms it IN WRITING afterwards. Most systems model that as an order with a flag. This one models it as the law reads: the call's output is an OFFER — a confirmation code, an expiry clock, and NOTHING else. No order. No service. No commission. The letter in the customer's inbox says the honest thing — "Nothing is ordered yet" — and the confirmation is an AUTHENTICATED act: the signed-in customer confirms their own offer, which is a stronger "in writing" than any bare link, and it made the security fall out for free. A stranger with a stolen code sees a 404. A re-click orders nothing twice. Silence expires on a clock, and the partner's pipeline shows offered, confirmed, expired — three honest states.

Identity is earned

The cold prospect closed the loop with the arc's best design decision. Someone who is not a customer yet cannot receive an inbox letter — so the code returns to the partner, whose dialer's own SMS is the only channel that exists for this person. And identity is EARNED, not invented: the prospect registers with the offered email — the registration is the identity proof — signs in, and confirms. The offer binds to them at that moment. Order, activation, commission: the same spine as the warm base, joined at the instant the law says an agreement started existing. Commission too is born at confirmation, not at the phone call — because a partner paid for calls optimizes calls, and a partner paid for confirmed agreements optimizes honesty.

Three partner channels — the retail chain, the POS integration, the call center — on one attribution-and-commission spine, each shaped by its own regulation: the cash-register law stays in the certified till, the withdrawal law lives in the confirm-after-call. A channel is a set of legal facts with a UI.

Advice with receipts

"Is it possible to add some feature on the product side that recommends enhancements... can also do market analysis and suggest new products and prices?"

— the operator, asking for the most dangerous feature in the book

Dangerous, because an AI that suggests products and prices is one careless design away from an AI that invents them. The Product Advisor was built on the opposite premise: it COUNTS before it speaks. The top-up attach finding is arithmetic — customers holding both a plan and a top-up in the product inventory are the receipt that the plan's allowance is priced below their appetite — and the finding carries its numbers ("2 of 2 customers on this plan also bought top-ups") next to a ready proposal. The market half reads a per-tenant FEED SEAM, the same pattern as the banks and the ESPs: a mock tariff-comparison subscription in the demo, a real market-intelligence provider in production, the same hole. A rival meaningfully cheaper on the same data bucket becomes a finding with both prices on it. The tenant's LLM gets exactly one job: narrate what the arithmetic found. The numbers never come from the model.

The advisor that refused to find

The build's best moment was a failure that wasn't. The first mock rival sold ten gigabytes at 14.50 against our 15.00 — and the advisor said NOTHING, because a three-percent gap is inside the ten-percent materiality threshold, and an advisor that cries about noise teaches its owner to stop reading it. The mock was sharpened to 11.90 and the finding fired with the numbers on it. Two more receipts came out of the build: top-ups never appear as billing rates (they are one-time — the honest source is the inventory), and TMF lists cap at a hundred rows with the newest on the last page, so an advisor that doesn't paginate is an advisor that can't see recent history.

And every suggestion is a PROPOSAL, never an action: adopting a finding births a draft offering — lifecycleStatus "In study" — visible on the product owner's console, provably absent from the shop until a human promotes it. The Copilot's governance pattern, reused: the machine drafts, the person decides, the suite checks the draft did NOT leak into the active catalog.

An advisor earns trust the way an accountant does: by showing the arithmetic, by refusing to speak when the numbers are noise, and by never having the pen that signs.

The right-sized mind

"Is it possible to use different AI models for different tasks AT THE SAME TIME... summarizing and email writing with a cheaper model, tough tasks by a high-end model?"

— the operator, describing AI economics in one sentence

Summarizing an email and designing a product are not the same kind of thinking, and should not cost the same. The answer took one enum and a cache key — and THAT is the chapter's real lesson. Every AI feature in the system already spoke through one seam, the per-tenant LLM router; no call site had ever talked to a model directly. So "different models simultaneously" was not a rewrite: the router learned to resolve (tenant, tier) instead of (tenant), and the fleet of features sorted itself.

The task class lives at the call site

The design decision that matters is WHERE the classification lives. Not guessed from the prompt — declared in the code, by the developer who knows the stakes: campaign copywriting, knowledge answers, the CSR's 360 summary and the advisor's narrative are FAST; the product copilot's proposals and the next-best-offer judgment are SMART. Unclassified work defaults to SMART, because "when in doubt, spend" is the only safe default for a machine that speaks to customers and product owners. And the tiers go all the way down: not just the model but the WHOLE PROVIDER resolves per tier — provider, endpoint, key, model, each value tier-first, shared-second — so a local cheap endpoint can do the volume work while a frontier API does the judgment, for the same tenant, in the same minute.

The proof is on the wire, because routing that is merely configured is routing that is assumed. The mock LLM names the model in every answer and keeps a ledger of which API dialect served which prompt; the suite reads copywriting answered by the cheap model over the OpenAI dialect and the copilot by the frontier model over the Anthropic dialect — one tenant, two providers, two models, at once — while the tierless tenant on the stub is untouched, because tiering is per-tenant data, not a global rewire.

The seam discipline paid its dividend here: a feature that sounds like an architecture project — "multiple AI providers per tenant per task" — cost an enum, because for two years nothing was ever allowed to call a model directly.

Three homes for the bytes

"Before going ahead with S3, I also want to know if this will run on Azure."

— the operator, asking the question that saved the design

The product media had always lived in the honest-but-humble place: a BYTEA column in Postgres, with a comment in the migration admitting the ceiling. But when the object-storage work was finally scheduled, the codebase had a surprise waiting — past-us had already built the seam. A

`ContentStore` interface, an in-row default, and a comment predicting this exact chapter: "pointing an operator's S3 at the same API is one adapter class." The architectural decision had been made months earlier; only the adapters were missing. That is what a seam discipline buys: future features arrive as fillings, not surgeries.

The question that shaped the build

The operator's gating question — will it run on Azure? — turned a one-adapter task into the right two-adapter design, because the honest answer was NO, not automatically. AWS, MinIO, Cloudflare R2 and Google's interop mode all speak the S3 protocol; Azure Blob speaks its own dialect, SharedKey signatures and all, and the once-standard gateway shims are discontinued. A BSS whose front page says "vendor-neutral" and whose Terraform ships an AKS stack cannot treat Azure as an S3 pretender. So the seam grew two thin adapters — the whole SigV4 ceremony and the whole SharedKey ceremony, each in ~150 lines of plain HTTP, no SDKs — and the storage choice became one environment variable. The build's one bug was a spec detail worth keeping: against Azurite, Microsoft's own emulator, the canonicalized resource carries the account name TWICE — path-style addressing puts it in the URL, the signing rules demand it again in front — by the spec, not by accident.

The standard agreed all along

TM Forum, it turns out, had blessed the move before it was made. The Attachment pattern that runs through every TMF API is url-OR-inline by design — an attachment is canonically a reference to content that lives elsewhere, with inline bytes as the convenience case. And TMF667 is a document MANAGEMENT api: metadata, categories, lifecycle, over content whose physical home was always meant to be someone else's problem. Fronting Postgres BYTEA was the unusual arrangement; fronting an object store is the typical one. The suite proved it the house way: the same upload-and-serve loop against all three stores, with the DATABASE giving the receipt each time — in-row shows the bytes in the row; the cloud stores show content IS NULL beside an s3: or azure: key. The bytes provably moved; the API and every attachment URL never changed by a byte.

"Any cloud" was a claim on the README for months. Now it is a database row that says NULL where the bytes used to be — in two dialects.

The rungs before the ladder

"When to upgrade search with Elastic?"

— the operator, asking the question whose best answer is a date far in the future

Some upgrades are won by refusing them. The search question had an obvious expensive answer — an OpenSearch cluster, RAM-hungry, operated forever — and a boring correct one: Postgres has rungs, and the system had climbed none of them. So the decision became a framework instead of a purchase. Elastic earns its place when someone asks for the OMNIBOX — one field searching parties, tickets, orders and articles together, ranked — or when search traffic deserves its own scaling axis. Until then, the rungs: and when that day does come, the event backbone is already the indexing pipeline, because every party, ticket and article publishes to Kafka — the hard part of an Elastic adoption, built years early for other reasons.

A net that only speaks when strict is silent

The first rung was the typo net. The CSR search — multi-word, term-AND, name-ranked, the shape a real bug report once forged — stayed byte-identical BY CONSTRUCTION: trigram similarity only runs when the strict search returns nothing. A query that finds people answers exactly as yesterday; "solviæg fjelheim", two misspellings deep, now finds Solveig instead of an empty list. And it fails open — on a database without the extension, the empty answer simply stands. The upgrade's safety was not tested into it; it was designed into the control flow, and then tested anyway: the suite proves the four behaviors, and seven regression suites — including the dealer and telesales channels, which look customers up through the same path — prove nothing else moved.

Each tenant searches in its own tongue

The second rung fixed a bug that had been invisible in English. The knowledge base's full-text search was hardcoded to the English stemmer — so nova's Norwegian articles were being stemmed by rules for the wrong language, wrong by construction, unnoticed because the demos searched in English. The tenant's locale now picks the stemmer, with an index per language: "regning" finds "regningene" and "regn timer", which no amount of English morphology could do. Multitenancy turned out to reach all the way down into linguistics — the registry field that picks the storefront's currency now also picks how words break.

Meaning finds what words cannot

The third rung came days later, and kept the family shape: a net that only speaks when the layer above is silent. pgvector put an embedding beside every article — a vector column the JPA entity never maps, so the unit tests

never knew — and when full-text finds nothing, cosine neighbours answer under a strict ceiling. The suite's proof pair says everything: "why is my internet so slow" finds the fair-use article that contains NEITHER word, because throttling and slowness are neighbours in meaning-space — and "purple elephant orchestra" finds nothing at all, because a net that invents relevance teaches its readers to ignore it. The embeddings themselves are a seam, of course: a deterministic keyless stub (hashed words folded through a curated telecom synonym table, English and Norwegian both) for CI and air-gapped demos, any OpenAI-compatible model when an operator wants learned semantics. Three nets now stand in a row — trigram under strict, stemmer per tongue, meaning under keywords — each honest about when it speaks, and honest about when it stays silent.

Cost of the three rungs: one image tag, two trusted extensions, four indexes, zero new containers. Cost of the road not taken: a cluster, forever. The best architecture decisions are sometimes the ones that end with Postgres doing one more thing.

The boring engineering

"Do you think our BSS is production ready? Can it scale, does it have redundancy and failover — is it secured?"

— the operator, asking the only question that matters after fifty-five suites

The honest answer was uncomfortable and useful: the architecture was right and the operations were not. Fifty-five suites proved the system CORRECT; none of them proved it SURVIVABLE. And the first thing that would have failed in production was not exotic — it was arithmetic. Nine scheduled ticks move money and send mail: dunning notices, paper bills to the print house, campaign sends, commission hardening. Run any of those services twice — the very first thing a Kubernetes deployment does — and every tick fires twice. Two dunning letters. Two paper bills. The feature that makes a system scalable is, embarrassingly often, the feature that stops it doing everything twice.

The lock is a row

The fix kept the house style: no coordinator, no new dependency. A `tick_lock` table in each service's own database and seventy lines that play the ShedLock algorithm straight — claim a lease with an atomic UPDATE-then-INSERT, release when done, expire on crash so a dead replica never wedges the fleet. The claim runs in its own transaction, so a lost race can never poison the tick's real work. And the proof was the receiver's, not the sender's: the suite raised a SECOND billing replica against the same database — two schedulers racing every three seconds — cut a bill, and asked the mock print house how many copies it held. One. The twenty-seven outbox relays were left deliberately unguarded, and the reason is the point: at-least-once delivery with dedup at every consumer was designed in from the first event, so a duplicate relay is waste, never harm. Where the contract already absorbs duplication you leave it alone; where it moves money, you lock the row.

A ceiling on every door, and a backup that is not a hope

The dealer surface had rate limits because partners get retry storms; every other door was unmetered. Now the gateway runs two rings — the strict per-partner buckets it always had, and a wide generous ceiling on everything else, counted per person, per machine, per address, so one caller's flood never becomes everyone's outage. And the backup learned the book's oldest lesson: nothing exists until a suite has seen it. `backup.sh` dumps the fleet; `restore-drill.sh` restores it into a THROWAWAY container and goes looking for a sentinel customer created minutes earlier. The suite found her in the copy. An untested backup is a hope. This one is a fact with a timestamp.

What remains is written down where claims go to be audited — docs/hardening.md: managed HA databases, TLS inside the cluster, load tests, the compliance ledger. Production-ready is not a feeling; it is a list, with the unticked boxes showing.

The run that survives its own death

"go ahead with P1 hardening"

— the operator, four words, the morning after P0 shipped

The billing run had a secret it shared with most young systems: it was one transaction. Every customer's bill, one commit, all or nothing — which reads as integrity until you do the arithmetic. At a thousand customers it is a long transaction; at a million it is a transaction that never commits, and one bad account rolls back everyone's bill at minute fifty-nine. The fix inverted the shape: the GATHER phase reads, then each account's bill commits alone. And the resume logic cost nothing, because it already existed — the idempotency check that stopped double-billing a customer within a period became, unchanged, the checkpoint of a crashed run.

The suite's FIRST run failed — and the failure was the best find of the arc. The ledger showed what no code review had: a thousand and seventeen accounts accumulated by months of suite runs, a paced pass that outlived the client's patience, and the retry loop quietly starting a SECOND run while the first one's orphaned thread kept walking — two runs racing the same period, saved from double-billing by timing alone. Luck is not a guarantee, so two fixes followed the diagnosis: the run now takes the same row-lease the ticks take — a second trigger gets an honest *busy*, and the lease HEARTBEATS per account so a crashed run frees it in two minutes, not a run-length — and the database itself gained a unique index: one bill per account per period, as a constraint rather than a convention. Then the suite passed the way it was meant to: twelve customers, a kill mid-flight, a resume to exactly twelve bills, a ledger showing one run *superseded* beside one *completed*. The run has a face now; even its deaths are rows.

The ceiling that outlives its gateway

Po's rate limits had a caveat written beside them the day they shipped: the buckets lived in memory, so N gateways meant $N \times$ the ceiling and every restart was an amnesty. P1 moved the buckets behind a seam — in-memory for a laptop, Redis when replicas must agree — and the proof was a resurrection: trip the ceiling, kill the gateway, and the FRESH process refused the very first knock. A window that survives its enforcer is a window replicas can share.

Numbers, at last — with their caveats stapled on

Fifty-six suites had proven the system correct and none had ever timed it. A plain-Node load driver fixed that in an afternoon: on one laptop carrying all thirty JVMs, catalog browse answered six hundred and seventy-eight times a second at thirty-nine milliseconds p95; authenticated bill reads — every request paying the full JWT-

plus-row-level-security toll — four hundred and fifty-one. Floor numbers, the baselines document says plainly, not capacity claims; their job is to make regressions loud, and the suite's smoke SLO stands tripwire duty at ten times the baseline. The load test's first victim, satisfyingly, was the wide ring itself — the harness got throttled by Po's ceiling and had to ask for it to be lifted. The guardrails guard even the auditors.

And alerts exist now — ServiceDown, a stalled outbox, a 5xx front door, evaluated over thirty-two scrape targets where five had been. A dashboard nobody watches is documentation; an alert is a promise that silence means something.

The right to be forgotten, with receipts

"go ahead with P2 hardening"

— the operator, and the last unticked boxes were the ones the law owns

A European subscriber holds rights the way they hold a phone number: legally, specifically, and with someone obliged to answer. The right of access got the build's most satisfying privilege design: the data passport rides the CALLER'S OWN TOKEN. When a customer exports themselves, party-account walks every service that holds a piece of them — interactions, messages, marketing touches, carts, tickets, appointments, behavioural profiles — presenting the customer's own credentials at each door. Every service returns exactly what that token could already read, which means the right of access required no new authority anywhere: it is what the person could always see, finally gathered into one portable document. The DPO — the user-administrator wearing a second hat, honestly documented — may name another party. Nobody else can, and the refusal is a 404, never a 403.

Erasure that knows the law both ways

The eraser refuses twice, and both refusals are the feature. A subscriber with an ACTIVE plan cannot be erased — 409, terminate first, because erasure never breaks a running contract; and the check FAILS CLOSED, refusing when it cannot prove the inventory empty rather than guessing. And the bookkeeping categories — bills, payments, orders, usage, agreements — are retained with their legal basis NAMED in the report, because a tax authority's five-year claim on an invoice outranks a deletion request, and pretending otherwise would be theatre. What the law does not hold, goes: seven services delete their slices, the profile anonymizes IN PLACE — "Erased", no contact media, the id surviving as the pseudonymous peg the retained records hang on — and the login dies at the identity provider, name and email scrubbed from the realm. The report of all of it becomes an immutable audit row. The suite's favourite assertion is the quietest: after erasure, the person's own password simply stops working.

Clocks, card numbers, and honest boundaries

Retention became a clock — interaction records and settled tickets deleted on schedule by the same guarded ticks that stopped double-firing in Po; open tickets are never touched, and the dev default is "keep forever" so no suite loses its timeline. The PCI question got the verify-then-claim treatment: the payment vault's columns are token, brand, last-four, expiry — no card number has anywhere to live — an SAQ-A-shaped posture claimed only as far as a README may claim it, a QSA attesting the rest. And lawful intercept was answered the NRDB

way: shaped, not connected — the BSS half is warrant-gated disclosure, and it ships when an operator with a legal obligation names its national format.

Fifty-eight suites. The last one proves a person can leave — completely, audibly, and with the paperwork the law demands on BOTH sides of the deletion. A system that can forget you on purpose is a system that was honest about remembering you.

The TM Forum vocabulary, plain-English

Every component in this book implements a TM Forum Open API — a published REST standard for one slice of an operator's business. The full plain-English glossary of what each one *means* — what the standard is, what it's for, and how genalpha-bss uses it — lives alongside this book as a companion reference, grouped to mirror the acts you've just read.

→ [Read the TM Forum reference \(tmf-reference.md\)](#)

Three words to carry out of it: **TM Forum**, the industry's standards body, which publishes the Open APIs; **ODA**, its blueprint for building a BSS as composable components rather than a monolith; and **CTK**, the conformance kit that proves a component is the real standard and not a lookalike. This whole book is what those three words look like when someone actually builds them.