

A BUILD MEMOIR

Verify Everything

*How one person and an AI built a
complete telecom BSS — one proof at a
time*

Built with **Claude Fable 5** · genalpha-bss · thirty-one components, six channels, and a discipline

What this book is

This is the story of building a telecom Business Support System — the software an operator sells, bills, and serves customers through — from an empty repository to thirty-one composable components, six customer surfaces, and twenty-one end-to-end test suites that all pass in a real browser. It was built by one person working with an AI pair, over a compressed and relentless arc, held together by a single discipline: *nothing is real until it is proven, end to end, against the real thing.*

It is a *memoir*, not a manual. There is a reference book to be written about how the pieces fit; this is not it. This is about the decisions, the dead ends, the two-in-the-morning debugging, and the method — and about what it means that a system this size was built the way it was.

Three convictions run through every chapter. The first is **vendor-neutrality**: any identity provider, any database, any message broker, any AI model, any payment processor, any number-porting authority. Nothing operator-specific is hardcoded, ever. The second is **verification**: official conformance kits, real-database migration tests, isolation proofs, and every feature driven through a real browser. The third is **honesty about what's a stand-in**: a thin network layer, a mock payment processor, a shaped-but-unconnected clearing house — each named plainly, so a demo stays credible instead of turning naive.

And beneath all three, the thread that makes this book unusual: it was built with an AI as the executing partner — Anthropic's **Claude Fable 5**, the most capable generally-available model of its moment. Not autocomplete. A collaborator that wrote the services, argued the architecture, made real mistakes, and got caught by the tests. The book is, quietly, a field report on what building at this scale with a model this capable actually feels like.

On the jargon. This is a book about a TM Forum ODA BSS, so it can't avoid the vocabulary. The rule is that it never assumes you already know a term: each chapter explains the one or two standards it needs as the story reaches them, and the appendix collects the full "what every TMF component means" reference. You should be able to read this having never heard of TM Forum and come out fluent.

Every claim in these pages has a commit, a test, or a screenshot behind it. That is not a stylistic choice. It is the whole argument.

ACT ONE

The First Cut

*Why build a Business Support System at all — and
the discipline that would define everything that
came after.*

The lock-in that started it

"The most dangerous phrase in the language is: we've always done it this way."

— Grace Hopper

Every telecom operator on earth runs software to answer three questions: what can a customer buy, what do they owe, and how do we serve them when something breaks. That software is called the BSS — the Business Support System — and it is, almost universally, the place operators are most trapped. The catalog, the ordering, the billing, the care: bought as one enormous suite from one enormous vendor, welded together over a decade, impossible to leave. The industry has a word for the escape hatch it wishes it had — *composable* — and a standards body, TM Forum, that publishes the blueprint for it. What almost nobody has is the thing built end to end, out in the open, where you can watch it work.

So that was the bet. Build the composable BSS the industry keeps promising, as a real product, and make it *vendor-neutral* from the first line of code. The operator I know best runs one particular identity stack; the product must never depend on it. It must run against any OpenID Connect provider, any PostgreSQL, any Kafka-protocol broker. Neutrality wasn't a feature to add later. It was the shape of the thing.

The first five components were the spine of any BSS. A **product catalog** (the shelf: what exists to sell, at what price, in what bundles). **Product ordering** (the checkout and its aftermath). **Product inventory** (what each customer actually has). A **party** component (the people and companies you do business with, kept deliberately separate from their login). And a **gateway** — one front door, so the outside world touches exactly one thing. Each spoke a published TM Forum Open API, which meant that from day one the system talked in a language other vendors' software already understood.

Constraints chosen early are cheap on day one and priceless on day one hundred.

That is the whole thesis of the first act, and it would be tested again and again. Every time the temptation appeared to hardcode something — the IdP, the number format, the country — the neutrality constraint said no, build a seam. Seams cost a little now. But a system made of seams is a system you can take apart, swap pieces of, and sell to an operator who already owns half of it. A monolith is a system you can only sell whole, to someone who owns none of it. There are not many of those customers left.

The lesson. The expensive decisions in software are the early, invisible ones. Choosing "any IdP, any database, any broker" on the first day looks like over-engineering. It is the opposite: it is the only cheap moment to make that choice.

Proof, or it didn't happen

"Simplicity is prerequisite for reliability."

— Edsger W. Dijkstra

The second decision mattered more than the first, though it looked like a chore. How do you *know* a thing works? The easy answer — write a unit test, watch it go green — is a comfortable lie. A unit test proves that the code does what the code's author thought it should. It says nothing about whether the author understood the problem, whether the pieces fit, or whether a real customer clicking a real button gets a real result.

So the answer became doctrine, and the doctrine was severe. First: the original components would pass TM Forum's official **Conformance Test Kits** — an independent, published suite that proves a component really implements a standard rather than merely resembling it. Not "we think it's TMF620." An outside authority agrees it's TMF620. Five kits at first; the doctrine would compound until, near the end of this story, *eleven* kits agreed — and the two that deliberately didn't became one of the book's most useful lessons. Second: every feature, without exception, would be driven through a *real browser* — register, configure, add to cart, pay, watch the order complete, watch the notification arrive. If a human couldn't do it in a browser, it wasn't done.

There was a subtler third rule, learned the hard way. The tests ran against an in-memory database that imitated PostgreSQL. Close, but not identical — and "close but not identical" is precisely where the migrations break. So the migrations grew their own test that spun up a *real* PostgreSQL and ran every schema change against it. The in-memory database is for speed. The real one is for truth.

Verification isn't overhead. It is the thing that lets you move fast and still trust the result.

This is the discipline that would, later, make it rational to build alongside an AI at a speed no solo human could sustain. When the model wrote a service, the tests didn't care that it was confident. They cared whether an order actually placed, whether a row was actually isolated, whether a real browser actually saw the result. The first "ALL CHECKS PASSED" scroll past the terminal was not a milestone. It was the establishment of a court of law that would preside over everything that followed.

The lesson. Choose your arbiter before you need it. A test that drives the real system through the real interface is a fact. Everything else — the author's confidence, the model's fluency, the green unit test — is an opinion.

The pair

"Programs must be written for people to read, and only incidentally for machines to execute."

— Harold Abelson & Gerald Jay Sussman

The collaborator has a name: **Claude Fable 5**. It wrote most of the lines in this system. Describing the working method is the honest heart of the book, because the method is the thing that's actually new. A human held intent and judgment — what to build, whether it was enough, which architectural fork to take. Fable 5 held execution and breadth — the services, the boilerplate, the cross-component wiring, and the ability to keep an entire thirty-component system in view at once, which no human can.

What it felt like, most of the time, was uncanny. You describe a component — "a promotion service, anonymous code validation, redemption tied to an order, a discount line on the bill" — and a few minutes later it exists, wired into the gateway, event-publishing, tenant-isolated, with a test that drives it through the browser. The model doesn't just write the service; it remembers the seventeen other services it has to speak to, and the tenant rules, and the event envelope, and it wires them without being told.

But a memoir that only showed the magic would be the same lie as a passing unit test. Fable 5 was wrong, sometimes, and wrong in the specific way capable models are wrong: fluently. It misdiagnosed a gnarly network bug twice before the evidence forced the truth. It once added a configuration value to the wrong service because a text-replacement matched the first occurrence instead of the intended one — a mistake no human would make and no human would catch by reading, but the test caught it in seconds, because the number never activated on the ported line.

The verification discipline isn't a brake on a capable model. It is the thing that lets you use the capability at full speed without fooling yourself.

That is the shape of the partnership, and it will recur in every chapter that follows: the model proposes at a breadth and speed that is genuinely startling; the human decides and, crucially, the *tests decide too* — not the model's confidence, and not the human's. Every commit in the repository is co-authored by Claude Fable 5. The credit is on the record, not just in the prose.

The lesson. AI-assisted building is only as trustworthy as the verification around it. Give the best model the whole system to hold, and a discipline that catches it when it's wrong, and the pair can build what used to take a team — provided nobody ever mistakes fluency for correctness.

ACT TWO

Scale & the Multitenant Reckoning

*Five components became thirty. One deployment
learned to hold two operators who must never see
each other. And the whole invisible machine learned
to show its face.*

One by one


"A distributed system is one in which the failure of a computer you didn't know existed can render your own computer unusable."

— Leslie Lamport

After the first five, the components came like tide. Stock, so the shop could tell the truth about the last handset. Payment, so money could move. Billing, so a customer could be charged for what they used. Qualification, so an address could be asked whether fibre even reached it. Appointments, tickets, interactions, communications, carts, usage, agreements, promotions, roles, addresses, recommendations, the card vault, documents. Each one a small, self-contained business, each one speaking a published standard, each one wired into the same nervous system.

That nervous system was the thing that made thirty components a system instead of thirty programs. The rule was simple and never broken: a component never reaches into another's database, and never calls another and waits. It writes what happened to its own tables, in the same transaction as its own work, and a background process ships that record onto a shared log. Other components listen. When an order completes, the ordering component doesn't *tell* billing, inventory, and communications to act. It simply records, in its own words, that an order completed — and the others, listening, do their part.

No component tells another what to do. It records what happened; the others listen.



The transactional outbox: write-what-happened, ship-it-to-the-log, let listeners react. Loose coupling as a physical property of the code, not a diagram.

Why go to this much trouble? Because the alternative — component A calls component B and waits — is the thing Lamport warns about. Chain enough synchronous calls together and a hiccup in a component you'd forgotten existed freezes a checkout. Record-and-react has the opposite failure mode: if billing is briefly down, the events wait patiently on the log, and billing catches up when it wakes. The customer's order completed the moment the order completed. Everything else is downstream, and downstream can be a little late without anyone noticing.

Thirty programs become a system when they stop calling each other and start telling the truth about what just happened.

The lesson. The coupling in a system is not in its diagram; it's in its call graph. Components that record facts and react to facts are loosely coupled in the only way that matters — the way that survives one of them falling over at 3 a.m.

Two operators, one deployment


Data isolation is not a feature you add. It is a property you can lose — silently, on any query, forever.

— from the build log

Here is the reckoning the second act is named for. A product you can sell to many operators is a product that runs many operators *at once*, on one deployment, sharing the same tables — and it must be impossible, not merely unlikely, for one operator to see another's customers. Get this wrong and there is no apology that suffices. This is the chapter where the whole architecture had to earn its trust.

The first question was: what is a tenant? The tempting answer is "a row in a tenants table," but that just moves the problem. The answer chosen was harder and truer: a tenant is a *verified identity provider*. If you can prove you were authenticated by the issuer that operator A registered, you are operator A's user — no separate account system, no shadow directory, no operator-specific code. Bring your own OpenID Connect issuer and you bring your own tenancy. That's neutrality and multitenancy solved by the same stroke.

But identity at the door is not isolation at the database. The load-bearing wall is *Postgres row-level security*. Every tenant-owned table carries a `tenant_id`, and the database itself — not the application, which can have bugs — refuses to return a row whose tenant doesn't match the one stamped on the current connection. The application runs as a restricted role that *cannot* see across tenants even if a query forgets its `WHERE` clause. There is a narrow, audited escape hatch for genuine system work, and migrations run as the owner. Everything else lives inside the wall.



Three layers, and only the last one is trusted. Identity resolves the tenant; the connection carries it; the database enforces it. A forgotten `WHERE` clause is caught by the wall, not by hope.

And then — because this is a book about proof — the isolation got its own test. Not a comment swearing it was safe. A test that logs in as operator A, creates a customer, logs in as operator B, and demands that B cannot see it — against a real PostgreSQL with the real policies. The gateway got its own trick too: an anonymous visitor typing a shop's hostname is routed to the right tenant by the hostname alone, so a customer who hasn't logged in yet still lands in the right operator's world.

A promise of isolation is worth exactly the test that tries to break it and fails.

The lesson. Put the guarantee at the lowest layer that can enforce it. Application code that "always filters by tenant" is one forgotten clause from a breach. A database that physically cannot return the wrong row is a different kind of promise — the kind you can sell.

War stories of scale

"Plan to throw one away; you will, anyhow."

— Frederick P. Brooks Jr.

Thirty services do not fit on a laptop by wishing. This is the chapter of the unglamorous fights — the ones that never appear in an architecture diagram and consume whole evenings. The container base image that had no build for the laptop's chip, so every image had to be pinned to one that did. The memory: thirty Java processes, each wanting a comfortable heap, summing to more RAM than the machine had, until every service was told firmly how little it was allowed. The builds parallelised until the fan roared, then throttled to four at a time so the machine stayed usable.

The test harness had its own private hell. It spun up real databases in throwaway containers — the right instinct, the one that catches migration bugs — but it talked to the container engine over a socket that, on this particular setup, lived somewhere non-standard, and it guessed an API version the engine didn't speak. Two environment variables, discovered over an hour that felt like three, and it worked. The kind of fix that is one line in the end and a small saga in the getting-there.

None of this is architecture. All of it is the tax you pay to run the architecture, and a memoir that skipped it would be selling the same fantasy as a demo that only ever runs on the presenter's machine. The reason to write these down is not nostalgia. It's that the next person — or the next model — hits the same walls, and a paragraph here saves them the evening.

The distance between "it works" and "it works on a machine that isn't yours" is where most software quietly dies.

The lesson. The environment is part of the system. A build that only succeeds on one laptop, with one chip, one socket path, one memory profile, is not done — it's performing. Write the walls down; they are load-bearing knowledge.


Making it visible

"Perfection is achieved, not when there is nothing more to add, but when there is nothing left to take away."

— Antoine de Saint-Exupéry


A customer never sees a component. They see a shop. For twenty-odd components, the system had been real but faceless — provably correct, and invisible. This is the chapter where it grew four faces: a **storefront** where a customer browses and buys; a **CSR console** where an agent serves them; an **admin console** — the back office — where an operator runs their catalog, campaigns, and staff; and a **mobile app** in the customer's pocket.

The subtle demand was that the faces be the operator's, not mine. A tenant is a brand, and a brand has a colour. So the channels learned to wear the tenant: the same storefront, rendered in one operator's identity and then another's, from nothing but their configured brand and logo. The two shops below are the *same code*, the same components, the same database — wearing two different operators.



One storefront, two operators. Same components, same deployment, different brand — white-labeling as a runtime property of the tenant, not a fork of the code.

The mobile app closed the loop: the same BSS, the same tenant rules, in the customer's hand.



The fourth channel — the account in your pocket, served by the same thirty components.

Making the system visible changed what it *was*. A correct-but-invisible platform is a thing you have to be an engineer to believe in. A shop you can click through, in an operator's own colours, is a thing you can show a buyer. The act closes with the machine no longer merely proven, but legible.

The lesson. The channel is where correctness meets belief. All the conformance in the world persuades no one who can't see it; a shop that wears the customer's brand and just works is the argument the tests were always making, finally visible.

ACT THREE

The Intelligence Turn

Adding AI to a BSS without surrendering to a single vendor, without leaking a customer's secrets, and without pretending a rule is a brain.

AI on our terms

"All models are wrong, but some are useful."

— George E. P. Box

The moment you add AI to a product, a gravity well opens: one vendor's model, one vendor's API, one vendor's prices, one vendor's outage becoming yours. For a platform whose entire conviction is neutrality, capitulating there would have been a betrayal of the first chapter. So the AI arrived the same way everything else had — behind a seam. An *intelligence* component with one job: to be the only place in the system that talks to a language model, and to talk to *any* of them.

A stub model for tests that need no cloud. An OpenAI-compatible adapter for the dozens of providers that speak that dialect. An Anthropic adapter for Claude. Each selected by a flag, per tenant, so operator A can run one model and operator B another, on the same deployment. The model is a setting, not an architecture.

Two rules were welded on before the first feature shipped, because they are the two ways AI quietly poisons a BSS. The first: **always-on redaction**. Before any customer text reaches any model, emails and phone numbers are stripped to `[email]` and `[phone]`. Not optionally. Not configurably. Always. The second: **an audit ledger**. Every prompt and every completion is written to a per-tenant, isolated table, so an operator can answer the question every regulator will eventually ask — *what did the AI see, and what did it say?*

A model you can swap is a model that can't own you. A prompt you can audit is a model you can defend.

This chapter is where the first act's conviction met the decade's most seductive lock-in and did not blink. Everything intelligent that follows — the copy assistant, the copilot, the churn engine — lives behind this seam, redacted and logged, portable across models by a single flag.

The lesson. Adopt the capability, refuse the capture. An AI seam with redaction and audit welded in is how you get the upside of a language model without betting your product, your customers' privacy, or your compliance story on one vendor's roadmap.

Where AI belongs

The right question is never "can AI do this?" It is "should a human be reading a draft, or reading a decision?"

— from the build log

A seam is a capability, not a use. The harder question was where in a BSS an AI actually earns its place — and the answer that held up was consistent: *AI drafts; humans decide*. It writes the first version; a person ships it. Every feature was chosen to sit on the draft side of that line.

The first was a campaign copy assistant. A marketer building a customer journey types a rough intent and gets back a subject line and a body, on-brand, with the promo code guaranteed to appear — but never inventing the discount amount, a bug caught early when a model cheerfully promised "20% off" for a code worth ten. The marketer edits and sends. The AI saved the blank page, not the judgment.

GenAlpha

BACK OFFICE

demo

Sign out

Product Offerings

Product Specifications

Product Offering Prices

Product Stock

Customer Bills

Serviceable Areas

Appointments

Campaigns

AI Audit

Campaigns

36 total

NEW

Name *

Summer win-back

Trigger *

Churn risk detected (AI scorer)

State filter

e.g. completed (optional)

Promo code to include

—

Message subject *

Message — {code} inserts the promo code *

AI assist

warm win-back for customers whose plan is ending, offer the code

Draft

Create

NAME	STATUS	TRIGGEREVENTTYPE	PROMOTIONCODE	REACHED	
Smoke welcome	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783787634183	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783787720216	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume
Welcome journey 1783789023110	paused	ProductOrderCreateEvent	WELCOME10	1 customer	Resume

The admin console's campaign builder. The AI drafts the subject and body; the marketer owns what actually goes out.

The second went to the support desk: a CSR copilot. When an agent opens a customer, the copilot reads the 360 — the products, the tickets, the timeline — and offers a two-line summary and a suggested reply. Crucially, it holds no credentials of its own; it sees exactly what the agent is allowed to see, through the same tenant wall, and not one row more. The agent reads a draft, not a verdict.

GenAlpha

CSR CONSOLE

GENALPHA-RETAIL

Customers

Tickets

Stock

AA Anna Agent

Sign out

GG

Gina Guest

013f3d14-b1ba-43a7-b978-2e0071df3f97 · Gästgatan 7, 22233 Göteborg

The provided customer data is summarized deterministically (stub provider — configure a real model for genuine insight). Nothing here left the machine.

Review the open items in the 360 with the customer.

Check whether the current plan still fits their usage.

Drafted by stub (deterministic-template) — verify before acting.

Orders

GenAlpha One Home & Mobile

Jul 11, 12:34 AM

acknowledged

Complete

Cancel

Services

Nothing provisioned.

Usage this month

No usage recorded.

Agreements

No agreements.

Bills

No bills.

Appointments

Installation: GenAlpha One Home & Mobile

Jul 14, 01:00 PM

cancelled

Active cart

The customer's cart, live — assisted checkout starts here.

Promotions & payment

No promotions or saved cards.

Suggest next

GenAlpha One Home & Mobile

#1

GenAlpha TV Max

#2

Samsung Galaxy S26

#3

Tickets

No tickets.

Raise a ticket for this customer...

Raise ticket

Interactions

No interactions logged.

The CSR copilot: a customer summary and a suggested reply, drawn only from what this agent may already see. It drafts; the agent decides.

The discipline is quiet but total. Nowhere does the AI place an order, issue a refund, or close a ticket on its own. It compresses, it suggests, it drafts. The moment of consequence stays human — which is exactly why an operator can turn it on without fear.

The lesson. The safe and valuable place for AI in a system of record is the draft, not the decision. Put it where a human reads its output before anything irreversible happens, and you get the speed without the liability.

The engine that learns

A rule you can read is honest about being a guess. A model that learns is honest about being wrong until it isn't.

— from the build log

Churn — a customer about to leave — is the prediction every operator wants and few trust, because most churn scores are black boxes that can't say *why*. So the churn engine was built to be honest twice over.

It began as rules a human can read and argue with: a commitment ending within the month, an allowance running consistently against its ceiling, a cluster of support tickets, a service degraded during an outage. Four plain signals, each defensible, each explainable to the customer it flags.

But rules are a floor, not a ceiling, and the operator's real ask was pointed: *"I want it to learn once it's live, and become production-quality in a few months — or be trained on old data."* So underneath the rules sits a real, if modest, learner. Every day it snapshots each customer's features; when customers actually leave — or actually port their number out, the strongest signal of all — it labels the outcome. From those snapshots and outcomes it fits a per-tenant model that sharpens over time. Cold-start on rules; warm up on reality; import history if you have it.

Start with a rule you can defend. Earn the right to a model you can trust.

And because it lives behind the same event backbone as everything else, a churn signal isn't a number on a dashboard nobody reads. When the engine flags a customer, it emits an event — and the campaign component, listening, can reach out before the customer reaches for the exit. The prediction becomes an action without a human relaying it. That closed loop — detect, decide, act — is the whole point, and it is built from the same record-and-react nervous system as an order completing.

The lesson. Don't ship a black box into a domain that demands explanations. Ship the readable rules first, let them earn trust, and grow the learned model underneath where the stakes reward it — with the outcomes reality hands you as the labels.

ACT FOUR

Autonomy Accelerated

*A stadium, a 5G slice, edge GPUs sold by the token,
and a network that heals itself while a B2B deal
closes itself. The Catalyst, rebuilt on our own BSS.*

The Catalyst

"The best way to predict the future is to invent it."

— Alan Kay

There was a proof-of-concept I'd seen at an industry event and helped write up: selling the network edge as an AI platform. A football stadium wants a private 5G slice with ultra-low latency and on-site AI capacity for real-time video. A business salesperson uses an AI assistant to shape the lead; the operator's systems autonomously judge whether the network can even do it, propose an edge-AI upsell, price it, provision it — and then, when a fibre is cut mid-event, heal around the damage and keep the promise. It was a vision demo, stitched from many vendors' parts.

The dare I set myself was to rebuild it — not as a slideshow, but running, on the very BSS this book has been building. Every glossy moment of that vision would have to become a real request against a real component that really did the thing. If the BSS was what I claimed, it should be able to *host* the future, not just narrate it.

A vision demo shows you the destination. The dare was to lay the track.

This act is that rebuild, moment by moment: intent, feasibility, the upsell, token pricing, self-healing, and — the strangest and most modern part — one company's AI agent negotiating with another's, machine to machine. Each one grounded in a component you've already met, extended just far enough to carry the weight.

The lesson. The test of a platform is whether someone else's flagship demo can be rebuilt on top of it as working software. A BSS that can host the industry's most ambitious vision, honestly, is a BSS that was built right.


Intent, feasibility, and the honest "no"

Tell the network what you want, not how to build it — and let it tell you the truth about physics.

— from the build log

The old way to sell a network service is a form with forty fields. The new way, the one the standards call *intent*, is a sentence: "a sub-20-millisecond slice at the stadium, with AI inference capacity." The customer states the *what*. The system works out the *how* — or works out that there is no how.

That last part is where most autonomy demos cheat, and where this one refused to. The orchestrator runs a feasibility check against a hard, physical rule: a sub-20ms guarantee requires an edge-GPU pool actually covering the place in question. If one exists, the intent is feasible and the system proposes the slice. If none does, the answer is **INFEASIBLE** — and the message says, in as many words, that this is *physics, not policy*. It will not promise latency that light itself can't deliver.



Autonomous feasibility with a spine. When the physics allows it, the system proposes — and upsells edge-AI. When it doesn't, it says so plainly, and doesn't sell a lie.

When it is feasible, the same autonomy that could have said no instead says "yes, and": a stadium buying low latency for video is a stadium that wants AI inference at the edge, so the system proposes it unprompted. The upsell isn't a hardcoded banner; it's a consequence the orchestrator reasons its way to. Autonomy that can refuse is autonomy you can believe when it recommends.

The lesson. An autonomous system that can only say yes is a liar with a schedule. The credibility of every "yes, and here's the upsell" rests entirely on its willingness to say "no, and here's why" when the physics demands it.

Priced by the token

The unit of billing tells you what business you're really in.

— from the build log

Here the BSS earned its keep in a way no slideshow can. An edge-AI service isn't sold by the month; it's sold by the inference — by the *token*, the same unit the AI industry itself runs on. So the catalog gained an Edge AI Inferencing offering with real token metering: fifty million inferences included, then a per-million overage price, defined the standard way and rated the standard way.

This is where the earlier, unglamorous components suddenly paid off. Usage metering — built chapters ago for gigabytes and minutes — didn't care that the unit was now AI tokens. Billing didn't care. The commitment, the overage, the line on the invoice: all of it already existed, because it had been built as a general capability rather than a special case. Selling AI by the token required almost no new billing code, because the billing had been built right the first time.

The best feature is the one you already built for another reason and discover you can charge for.

And the commercial motion around it was the intent loop made concrete: the intent produced a priced **quote** — born from the conversation, itemising the slice and the token bundle, wrapped in an AI-written narrative explaining what the customer was buying and why. Accept the quote and it became an order. Lead to intent to quote to order, most of it autonomous, all of it against real components.

The lesson. Build capabilities, not features. A usage-and-billing engine built generically will one day meter a unit you hadn't imagined — AI tokens — and charge for it without a rewrite. That optionality is the compounding return on doing it properly.

The network that heals itself

Resilience is not the absence of failure. It is the presence of a plan that runs without you.

— from the build log

The dramatic beat of the original vision was a fibre cut, mid-event, and a network that healed around it before anyone in the stadium noticed. It's the moment that separates a brochure from a system, because healing is easy to *say* and unforgiving to *do*. So it got built as an actual loop, wired to the assurance component that had been watching the network all along.

A critical alarm arrives on the fibre route. Assurance, which turns alarms into tracked service problems, raises one. The self-heal service sees it and acts: it re-homes the affected service off the cut fibre and onto an edge GPU site that can still carry it, records the migration, opens a ticket so a human has an audit trail, and — once the service is stable on its new path — resolves the problem automatically. The SLA holds. The stadium keeps its slice. And every step left a record, because a self-healing system that heals invisibly is indistinguishable from one that's lying.



Every step recorded. A heal you can't audit is a claim, not a system.

The self-heal loop: alarm to problem to re-home to restored SLA, autonomous end to end, with a ticket left behind so a human can trace exactly what the machine did.

Autonomy you can trust is autonomy that shows its work.

The lesson. A self-healing system must be an *auditable* self-healing system. The re-home is the clever part; the ticket it opens on the way is the part that makes an operator willing to let it run unattended.

Agents talking to agents

The next integration boundary isn't an API. It's a conversation between two machines that each brought a model.

— from the build log

The strangest and most forward-looking beat came last. In the original vision, the buyer's side and the seller's side each had their own AI, and the deal moved between them machine-to-machine. That sounds like science fiction until you build it, and then it's just a protocol.

So the BSS grew a small server speaking the emerging standard for exactly this — a way for an *external* AI agent to drive the lead-to-order motion through a handful of well-defined tools: draft an intent, submit it, create a quote, accept a quote, list quotes. A salesperson's Claude, sitting entirely outside the BSS, could now shape a stadium deal by calling these tools — and on the other side, the operator's autonomous orchestrator answered with feasibility, an upsell, and a price. Two AIs, two companies, one deal, no human typing into a form.

What made this more than a party trick was that the tools weren't a special AI-only backdoor. They were thin wrappers over the exact same components a human channel used — the same intent, the same quote, the same order, the same tenant wall. The agent got no special powers and no special access. It simply drove the front door with its hands instead of a mouse.

When the integration layer becomes a conversation, the thing you must get right isn't the AI. It's that the AI touches nothing a human couldn't.

The lesson. Machine-to-machine autonomy is coming through the tool interface, and the discipline that makes it safe is the oldest one in this book: the agent acts through the same guarded front door as everyone else, with the same tenancy, the same limits, the same audit. No backdoors for robots.

ACT FIVE

Making It Real, Making It Legible

Hardening the money and the identity, keeping and leaving a phone number, running the whole thing on a real cluster, and finally making the invisible machine tell its own story.

The money and the identity

Two things a telecom must never get wrong: charging a customer twice, and believing someone is who they aren't.

— from the build log

"Let's harden the BSS first" was the instruction, and it was the right one. A demo can be forgiven a hand-wave at payment. A product cannot. So the payment path was rebuilt to the standard a real processor demands. **Idempotency** first: every payment carries a correlator, so a customer who double-clicks, or a network that retries, gets the *same* payment — never a second charge. Then the full life of money: authorize a hold, capture it, void or refund it, each a real movement recorded with the processor's own reference.

And because real payments sometimes demand more — a bank insisting on a second factor — the path learned to pause. When the processor requires strong customer authentication, the API doesn't fail; it returns a "we need more" with the URL to complete it, and resumes when the customer has. The mock processor even carries test cards that decline or demand that step-up, so the hard paths get exercised on every run, not just the happy one. A real Stripe adapter sits behind the same seam, one flag away — because of course it does.

Identity got the same seriousness. Some purchases — a high-value plan, a regulated product — should require a *verified* identity, the kind a national eID like BankID provides. So the catalog can mark an offering as requiring verified identity, and ordering enforces it: no verified identity claim in your token, no order — a clear refusal with a header telling the channel to send you to step up. The verification itself is the identity provider's job, which is exactly where a neutral platform wants it.

Hardening is where a demo becomes a product: the same flow, now unable to charge you twice or take your word for who you are.

The lesson. The unglamorous guarantees — idempotent payments, real authorize/capture/refund, verified identity where the stakes demand it — are the line between something that demos and something you can put real customers and real money through. Harden before you scale.


Keep your number, or leave with it

The truest test of an open platform is whether a customer can walk away — and take their number with them.

— from the build log

Number portability is the quiet backbone of telecom competition. A customer will only switch operators if they can keep their number, and regulators mandate it precisely because it keeps operators honest. Building it meant confronting a truth the first chapter predicted: portability is *country-specific*. In Norway, a national clearing house called NRDB, under the regulator Nkom, brokers every port. Elsewhere it's a different authority with different rules, different formats, different cutover windows.

So porting got the same treatment as everything else: a seam. A porting gateway with a country-aware implementation — NRDB for Norway, shaped exactly as the real thing but not wired to it, and honestly labelled so — and per-country rules for number format, cutover timing, and regulator. "Vendor-neutral" had quietly become "regulator-neutral." The same code ports a number in Norway, Sweden, Britain, or the United States by consulting the rules for the country, not by branching on a hardcoded assumption.



Both directions of the door. Port-in draws the customer's real number instead of a pool number and holds the line until cutover. Port-out ceases the service, releases the number — and feeds the churn engine the truest label it will ever get.

Keeping your number when you arrive meant the orchestrator had to wait: instead of grabbing a fresh number from the pool, it holds for the port-in's cutover and then carries the customer's own number. Leaving with it meant building something the BSS didn't yet have — a service *cease* flow — because a port-out is also a termination. And a termination, the churn engine noticed, is the single most honest churn label in existence: not a prediction of leaving, but the fact of it, fed straight back into the model that tries to see the next one coming.

The lesson. The features that let a customer leave are the features that prove a platform is open. Build the exit as carefully as the entrance — country-aware, honest about stand-ins — and the same event that loses a customer teaches the system to keep the next one.

Running it for real

"In theory there is no difference between theory and practice. In practice there is."

— attributed to Benjamin Brewster

Everything so far had run on a laptop, in a chain of containers. That proves the software works; it doesn't prove the software *deploys*. Between the two lies a canyon full of the things that never show up until you cross it — resource requests that were fine on a laptop and starve a scheduler, config that was an environment variable and now must be a mounted secret, a component that came up instantly and now waits on ten others.

So the whole system was packaged the way it would actually ship — a Helm chart for Kubernetes, with infrastructure defined as code for real clouds — and then *soaked*: brought up on a real cluster and left to run, watched, prodded, and made to prove it could stand on its own for more than the length of a demo. A soak isn't a test that passes or fails in a second. It's the long, boring, essential question — does it stay up? — asked of the parts that only misbehave over time.

"It runs on my machine" and "it runs" are separated by a canyon, and the canyon is where products are made.

The soak had its own runbook and its own honesty: which components are core and which can be scaled back to fit a modest cluster, what to watch, what "healthy" means over hours rather than seconds. It is the least cinematic chapter in the book and one of the most important, because it is the difference between something that impressed a room and something an operator could actually run.

The lesson. Deployment is a feature, and soak is its test. A system that has never run unattended on a real cluster for hours has not been proven to run at all — it has been proven to *demo*, which is a different and smaller claim.

The invisible, made visible


If a business process runs and no one can watch it, you have built a machine of faith.

— from the build log

All those events, flowing between all those components, were the true life of the system — and completely invisible. You had to be an engineer reading logs to believe any of it was happening. The last component set out to fix that: **Live Flow**, a face for the nervous system itself.


The first attempt was an honest engineer's instinct — a graph of components with lines lighting up as messages passed between them. It was correct and it was useless to anyone who wasn't me. The feedback was exact and it reframed everything: *show it like a business-process engine shows a process. Like a layman can understand — a customer ordering, the order component activating, an order-created event, captured by something.* Not a wiring diagram. A story.

So Live Flow was rebuilt as a set of business processes told in plain English — Order to Activation, Lead to Assurance, Churn to Retention, Port-in to Keep-Your-Number, Port-out to Leave — each a row of labelled stages that *light up in order* as the real events arrive, teal for a normal step, violet when AI touched it, rose when something went wrong, fading as they cool.



Live Flow's process view: the real event stream, narrated as business processes a non-engineer can read. The stages light in sequence as it actually happens.

Then the last ask: *can I click a box and see the detail? Can the colour change the moment an event happens?* Yes and yes. Each stage became clickable, opening the specific events behind it; each customer's journey could be followed as an individual thread through the process. The choreography lit as a customer really moved through it — the invisible machine, at last, telling its own story to anyone in the room.



Left: stages lit and cooling as events land. Right: a single customer's journey traced through the process. Click any stage to see the events beneath it.

A system that can narrate itself to a stranger is a system that finally understands what it's for.

The lesson. Observability for engineers is logs; observability for the business is a story. The same event stream, told as processes a layperson can follow, turns an invisible distributed system into something a customer, a buyer, or a regulator can simply watch.

Rules without red deploys, and the audit

In telecom, every new product brings new business rules. The question that decides whether you have a product or a project is: does a new rule mean new code?

— from the build log

The question arrived plainly, the way the best ones do: *"New products have new business rules. Is it new code every time, or is there a rules page?"* The honest answer, at that moment, was both. Prices, bundles, allowances, promo codes — all data, changeable in the console with no redeploy. But a genuinely new *kind* of rule — "max two SIMs per order," "these two products can't be bought together" — was a code change. There was no rules engine, and pretending otherwise would have broken the book's one law.

So the thirty-first component was born: **policy**. Rules authored as pure data — JSON-logic conditions an operator writes in the console — evaluated at order time by a small engine that can never execute code, only answer true or false. The proof was the point: author a "max 2 per order" rule through the API, watch an order for three get refused with the rule's own message, watch an order for two pass, then *disable the rule with a row* and watch the same order for three sail through. The rule appeared, enforced itself, and vanished — and the deployment never moved. The same engine then learned to answer a second question: not just "may they?" but "at what price?" — pricing rules that adjust a subtotal by segment, by cart, by verified identity, stacking like a customer would expect, landing on the bill as labelled adjustment lines. And the bundle model grew the formal cardinality the standard always offered: mandatory components, optional add-ons a customer may include, and "pick exactly N of M" choice groups enforced at order time.

A rule you can add, enforce, and retire without a deployment is the difference between selling a product and selling a promise of future engineering.

And then, because the machinery was all data and seams, came the reckoning the second chapter had promised: *run the official conformance kits against everything else*. The kits were museum pieces — Node-16-era runners that mangled URLs on modern machines — so the first day was spent building a harness that made their numbers trustworthy at all. Then, component by component: measure the real gap, reshape to the spec, redeploy, re-run, and re-run the *application's* own tests so conformance never broke commerce. Shopping cart. Party roles. Stock. Usage. The customer bill — nineteen thousand assertions on that kit alone. Consumption reports. Each driven to zero failures, each committed with its number.

The most valuable finding was the two components that *refused*. Payment would not accept the kit's empty, amount-less payment — because in this BSS, creating a payment is the authorization, idempotent by correlator, and no conformance score was worth re-opening the door to a double charge. Communication would not send a message with no recipient. Eleven kits green, two gaps kept on purpose and written down in plain sight — because a conformance page that admits "here we are stricter than the spec, and here is why" is worth more to a buyer than a wall of unexamined green.

The lesson. Configurability is a spectrum, and the honest move is to say where you are on it — then move the expensive end (new rule types, new prices) into data. And when an external standard disagrees with a guarantee you chose deliberately, the answer isn't to surrender the guarantee; it's to document the disagreement and let the buyer judge. Eleven zeros and two explained exceptions beat thirteen unexamined zeros.

ACT SIX

The Operator's Week

*Companies get a door of their own, the shop
becomes a home, products learn what they are —
and the whole machine starts speaking Norwegian.*

Everyone gets a door

"Do we have full B2B support in our BSS?" The honest answer was no — and honest answers are where features come from.

— from the build log

Every chapter so far had served one kind of customer: a person. But operators sell to companies, and a company is not a big person. It has members whose lines it pays for, an administrator who is not the operator's staff, and an accountant who wants *one* invoice, not forty. The audit took an afternoon and returned a familiar shape: the party model could hold organizations, nothing connected people to them, and nobody but the operator could manage anything.

The foundation went in the way foundations had gone in all along. Organizations gained a parent; individuals gained a membership; a new customer-side role — the business administrator — could act only inside their own organization, a boundary enforced by live lookups in the party service, never trusted from the request. Billing learned consolidation: members' charges regroup under the company, each line still attributed to the person who used it, and if the party service is unreachable the run fails open to per-person bills rather than billing nobody. A sixth surface appeared at `/biz` — the business console — where Bianca of Acme adds people, orders lines for them, and reads the company invoice.

Two truths surfaced in the proving. The first was operational: the consolidated bill came out empty because billing's machine account lacked one read permission — and the fix did nothing until the service restarted, because machine tokens are cached. That lesson ("grant, then restart, then believe") would pay for itself twice more before the week ended. The second was structural: Bianca's people existed as party records but could not sign in, because every channel keys a customer to their identity-provider subject. So invitation became a first-class feature: adding a person with an email mints a real login through the TMF672 seam — hardcoded to the customer role, so the power grants no escalation — and pins the party's id to the new token subject. Erik signs in and lands on *his* line.

Where should Erik land? The storefront was the obvious answer and the wrong one. "I think B2B my-page should be separate," came the verdict, and the channel split in two by role: the administrator sees the org console; a plain member sees a work line, usage, SIM care, and a note that the company pays — nothing to buy, no personal bills. One channel, two faces, the same pattern the role-scoped consoles had established.

Then money. A company that brings forty lines does not pay list price, and the machinery for that already existed: the pricing-rules engine. The billing run began passing the paying organization into the rules context — a field that simply does not exist for consumers, so a corporate deal can never leak to retail — and "Acme pays 20% under list" became a row of data: 25 becoming 20 in the company's price view and on the consolidated

invoice, to the cent. When the raw JSON condition drew the fair complaint that account managers should not write JSON, the console grew two plain-English rule kinds: pick the company from a dropdown, type the percent. The engine speaks JSON; the humans never have to.

The shop becomes a home

A customer doesn't own a list of rows. They own a plan, some lines, a SIM in a phone, and a bill they'd rather not think about.

— design note, My page

The critique arrived as five observations in one message: the page mixed the commercial with the technical, the plan-change menu offered a Samsung as if it were a tariff, bundles hid their contents, and where were the ten- and fifty-gigabyte plans, the one-time data top-ups? Behind all five sat one missing idea: the catalog had no notion of what an offering *is*. So it learned. Categories — mobile plans, broadband, devices, TV, top-ups — became persisted, first-class facts, and everything else followed from them.

The page rebuilt itself around the MyJio idea: it recomposes around what the customer holds. A card per line of business — the bundle with its components nested, every numbered line with its own SIM, the latest bill with a pay link, a discovery card fed by the recommendation engine. Plan change became a real TMF622 [modify](#) order: same service, same number, validated hard (your own product, an existing plan, like-for-like by category) and blocked while a commitment window is open. A bundle's broadband tier upgrades in place because bundles decompose into per-component products, while the commitment binds the bundle itself. The SIM became real — ICCID and PUK minted at activation, PUK revealed on request, PIN pushed over the air through a pluggable SIM-platform seam — and a five-gigabyte top-up became a purchase that genuinely grows this month's meter, delivered through the event stream and idempotent per order.

Complicated products got their showpiece: Family Max. Two fixed components, "family lines — pick one or two" of three plans, a configurable phone whose colour and storage ride the order, up to two streaming extras, a twelve-month term, a base price plus a one-time fee while every chosen option prices itself. The storefront configurator learned pick-N-of-M — checkboxes that cap themselves live, an add-to-cart that gates below the minimum — and the order API enforces the same limits, because the UI is convenience and the order is law. And when the fair question came — "does a product owner need to use JSON?" — the console grew a bundle composer: components marked included or optional, choice groups with their pick-ranges and defaults, terms and categories as plain fields, and a promise that anything it doesn't understand survives a save untouched.

Then the products that aren't lines at all. Netflix on the operator's bill is sold like any offering but fulfilled with the *partner*: the orchestrator now reads the catalog category to decide fulfilment — partner services mint an entitlement code through a partner seam, security products activate as resource-less features, insurance provisions nothing and simply bills. The proof was cinematic in the most literal way: reviewing stills of the demo film caught the top-up minting a phantom phone line with its own SIM, and the film's Norwegian take caught the partner machinery failing for the second tenant — one missing per-realm grant. Both fixed, both reproven, both on camera in the next cut.

Kroner and tongues

"In Norway it will require Norwegian, and NOK as currency."

— the requirement, in full

The audit was two sentences long. Currency: eighty percent there, because every amount in the TMF models had carried its unit from the first chapter — EUR appeared in code only as a fallback. Language: zero percent there. And so the fix followed the house's oldest rule: the tenant manifest, which already told every channel its brand name and colour, learned two more words — `locale` and `currency`.

The channels grew a translation layer with one deliberate trick: English strings *are* the keys, so an untranslated string falls through as itself and the English tenant risks nothing. Money formats through the platform's own internationalisation from each price's own unit — "299,00 kr/md." in Norway, the historical English format untouched because a dozen test suites assert it. Keycloak's built-in Norwegian switched on per realm. And Nova Telecom, the white-label tenant that had existed to prove isolation, became a real Norwegian operator: NOK prices, a tariff ladder, a top-up called Datapåfyll, and personas to walk it — Nils on Min side changing plans in Norwegian, keeping his number at 199 kroner.

B2B demanded the same parity and delivered the best find of the week: the business console had its identity provider hardcoded to the first tenant's realm. It could never have served a second operator — not mistranslated, but structurally monolingual. Wired into the manifest like every other channel, it mounted at Nova's own hostname; Birgit Berg of Fjellheim AS runs her people and lines in Norwegian and reads one consolidated invoice at 299 kroner. When the header still read "BEDRIFT," the correction came from the person who knows: Norwegian operators call it *Min bedrift*. Renamed. The whole B2B machine — boundaries, consolidation, invitations — worked for the second tenant with zero backend changes, which is the sentence the first chapter was always trying to earn.

The tenant-isolation suite now walks an entire purchase in Norwegian — Legg i handlekurven, Til kassen — and the demo film ends its Norwegian act with a plan change and a company invoice, in kroner, before Live Flow rolls the credits. One build. Every language. Every currency. The manifest did it, not a fork.

Proof against tomorrow

"Let me know if our BSS is PQC proof." Five words, and suddenly the audit was of mathematics itself.

— from the build log

Every audit so far had asked whether the system did what it claimed. This one asked whether its arithmetic would survive an adversary that does not exist yet. Post-quantum cryptography sounds like a problem for cryptographers, but for a BSS it is a problem of *inventory*: where, exactly, does cryptography live in this thing? The grep took minutes and returned the most reassuring answer available: almost nowhere. A secure random generator for PUKs and invitation passwords. A SHA-256 challenge in the login dance — hash-based, and hashes only lose half their margin to a quantum computer. And that was the entire list, because identity, transport and payment crypto had all been pushed behind seams in chapter one, back when "vendor-neutral" was a conviction and not yet a security posture.

The vulnerable pieces were exactly two, and neither was ours. The identity provider signs every access token with RSA — the algorithm Shor's algorithm eats first. And transport is classical TLS, which matters more than it sounds: an adversary can record encrypted traffic *today* and decrypt it the day their quantum machine boots. "Harvest now, decrypt later" is the one post-quantum risk with a present-tense clock, and it is entirely a deployment concern.

So the honest answer went back in two sentences — not PQC-proof, PQC-*ready* — and the follow-up question wrote the afternoon's work: so what's ours to do? Three things, and two of them shipped the same day. The Helm chart grew an edge: TLS 1.3 only, with controller values that switch the key exchange to a hybrid — a classical curve and ML-KEM-768 riding the same handshake, so clients that speak post-quantum get it and everyone else falls back gracefully. And the SIM PUKs, sitting in the database as an honest dev shortcut, got a vault: AES-256-GCM with the card's own ICCID bound into the ciphertext, so a stolen row cannot even be replayed onto another card. Old plaintext rows quietly upgrade themselves the next time their owner asks to see them. AES-256 is already comfortable against Grover; the fix was classical hygiene wearing a quantum justification.

The third item was the fleet itself: thirty services on a JVM that had never heard of ML-KEM. "So what's stopping the backlog?" Nothing, it turned out, except doing it — and doing it honestly meant meeting three of Spring Boot 3.5's edges in one afternoon: Flyway had modularised its databases, so twenty-six services booted to "Unsupported Database: PostgreSQL" until each carried the new module; the gateway's Spring Cloud train refused the new Boot on principle until it was bumped two generations; and one test dependency fell out of the managed BOM. The subtlest lesson wasn't in any pom: a single transient image failure aborted the parallel build and left a *mixed* fleet that still reported healthy — old containers serving happily under new build logs. Trust `java -version` inside the running container; never trust the tail of a build.

Then the definition of done: all thirty services healthy on Java 25 LTS, and the full twenty-one-suite sweep — every channel, every tenant, every language — green. The remaining migration, the IdP's RSA signatures becoming ML-DSA, waits on the OIDC ecosystem, and when it lands the services will not notice: they have never hardcoded an algorithm, they follow the issuer's published keys. That is the whole argument, and it doubles as the marketing copy precisely because every claim in it has a receipt: a BSS that embedded its own crypto would face a post-quantum rewrite. This one faced a checklist, and the boxes that were ours to tick are ticked.

C O D A

What it means

*One person, one AI, one discipline — and a system
that used to take a team.*

The load-bearing wall

"Make it work, make it right, make it fast."

— Kent Beck

Thirty-one components. Six customer surfaces, one of them wearing two faces by role. Twenty-one end-to-end suites, all green in a real browser. Eleven official conformance kits passing with zero failures — and two components that fail their kits on purpose, with the reasons written down. Business rules, prices and corporate deals that change with a row, never a redeploy. Isolation proven against a real database — and then proven again in Norwegian, in kroner, B2C and B2B, from one build. Plans that change without losing the number; SIMs with a PUK you can ask for; top-ups that genuinely grow the meter; Netflix fulfilled as a partner entitlement because the catalog category decides what a product is. A bundle composer so product owners never touch JSON, and an order API that enforces what the composer promised. A national number-portability flow, both directions. A quantum-safe edge and encrypted card secrets, with the one remaining post-quantum migration parked at a seam. A network that heals itself and leaves a receipt. And a film — re-shot after every feature — that keeps catching real bugs, because everything on camera is real.

Look back and the three convictions turn out to have been one conviction wearing three coats. **Vendor-neutrality** — any IdP, any database, any broker, any model, any processor, any porting authority — was possible only because every dependency was a seam, and every seam was verified with a real implementation behind it. **Honesty about mocks** was survivable only because the mocks were named and the seams were real, so the day a real Stripe or a real NRDB plugs in, nothing above it changes. And **verification** was the wall that held both up.

That wall is also the answer to the question this book is really about. It was built by one person and an AI — Claude Fable 5 — at a pace no solo human could sustain and no unverified pair could survive. The model's breadth was extraordinary; it held the whole system in view and wired thirty-one components without losing the thread. But breadth without an arbiter is just confident wrongness at scale. The tests were the arbiter. Not the model's fluency, not my intuition — a browser, a real database, an official kit, agreeing or refusing.

Give the best model the whole system to hold, and a discipline that catches it when it's wrong, and one person can build what used to take a team — if they never stop verifying.

That is the memoir's whole argument, and the reason its title is an instruction rather than a boast. The future of building software isn't "the AI writes the code." It's a human holding intent and judgment, a capable model

holding execution and breadth, and a ruthless verification discipline making the partnership trustworthy. A thirty-one-component, multi-tenant, standards-conformant BSS, built this way, is the existence proof. Every claim in these pages has a commit, a test, or a screenshot behind it.

Verify everything.

The TM Forum vocabulary, plain-English

Every component in this book implements a TM Forum Open API — a published REST standard for one slice of an operator's business. The full plain-English glossary of what each one *means* — what the standard is, what it's for, and how genalpha-bss uses it — lives alongside this book as a companion reference, grouped to mirror the acts you've just read.

→ [Read the TM Forum reference \(tmf-reference.md\)](#)

Three words to carry out of it: **TM Forum**, the industry's standards body, which publishes the Open APIs; **ODA**, its blueprint for building a BSS as composable components rather than a monolith; and **CTK**, the conformance kit that proves a component is the real standard and not a lookalike. This whole book is what those three words look like when someone actually builds them.